

# Visión Artificial VIAR25

## Lab 1: PyTorch Setup, Camera Calibration & Image Transformations

Prof. David Olivieri

Uvigo

September 8, 2025

# Lab 1 Overview

## Lab Goals

Transform mathematical theory from Lecture 1 into working code. By the end of this lab, you'll have implemented the core algorithms that power computer vision systems.

## What We'll Implement Today

- ① **Environment Setup:** PyTorch + OpenCV + visualization tools
- ② **Image I/O & Manipulation:** Loading, transforming, visualizing images
- ③ **Camera Calibration:** Zhang's method with checkerboard detection
- ④ **Homography Estimation:** DLT algorithm with RANSAC
- ⑤ **Lens Distortion Correction:** Brown-Conrady model implementation
- ⑥ **Color Space Transformations:** RGB  $\leftrightarrow$  HSV conversion

## Lab Structure

- **Demo Code:** Step-by-step implementation with explanations
- **Exercises:** Hands-on coding challenges to test understanding
- **Questions:** Theoretical connections to deepen comprehension
- **Extension Challenges:** Advanced implementations for extra credit

# Environment Setup & PyTorch Fundamentals

## Required Libraries

```
1 import torch
2 import torchvision
3 import torchvision.transforms as transforms
4 import cv2
5 import numpy as np
6 import matplotlib.pyplot as plt
7 from sklearn.linear_model import RANSACRegressor
8 import glob
9
```

## Exercise 1.1: PyTorch Tensor Operations

**Task:** Create a function that converts between NumPy arrays and PyTorch tensors

```
1 def numpy_to_torch(img_np):
2     """Convert HxWxC numpy array to CxHxW torch tensor"""
3     # TODO: Implement conversion
4     pass
5
6 def torch_to_numpy(img_torch):
7     """Convert CxHxW torch tensor to HxWxC numpy array"""
8     # TODO: Implement conversion
9     pass
10
```

## Questions:

- Why do PyTorch and OpenCV use different channel ordering?
- When would you use GPU tensors vs CPU tensors for image processing?

# Image I/O and Basic Transformations

## Demo: Image Loading and Visualization

```
1 def load_image(filepath):
2     """Load image using OpenCV and convert to RGB"""
3     img_bgr = cv2.imread(filepath)
4     img_rgb = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2RGB)
5     return img_rgb
6
7 def visualize_images(images, titles, figsize=(15, 5)):
8     """Display multiple images in a row"""
9     fig, axes = plt.subplots(1, len(images), figsize=figsize)
10    for i, (img, title) in enumerate(zip(images, titles)):
11        axes[i].imshow(img, cmap='gray' if len(img.shape)==2 else None)
12        axes[i].set_title(title)
13        axes[i].axis('off')
14    plt.tight_layout()
15    plt.show()
16
```

## Exercise 1.2: Geometric Transformations

**Task:** Implement the transformation hierarchy from Lecture 1

```
1 def apply_transformation(img, transform_matrix):
2     """Apply 3x3 transformation matrix to image"""
3     # TODO: Use cv2.warpPerspective
4     pass
5
6 def create_similarity_transform(scale, rotation, translation):
7     """Create similarity transformation matrix"""
8     # TODO: Implement using lecture equations
9     pass
10
```

# Corner Detection for Calibration

## Demo: Checkerboard Corner Detection

```
1 def detect_checkerboard_corners(img, pattern_size):
2     """Detect checkerboard corners with sub-pixel accuracy"""
3     gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
4
5     # Find corners
6     ret, corners = cv2.findChessboardCorners(gray, pattern_size, None)
7
8     if ret:
9         # Refine corners to sub-pixel accuracy
10        criteria = (cv2.TERM_CRITERIA_EPS + cv2.TERM_CRITERIA_MAX_ITER,
11                    30, 0.001)
12        corners = cv2.cornerSubPix(gray, corners, (11,11), (-1,-1), criteria)
13
14    return ret, corners
15
16 def visualize_corners(img, corners, pattern_size):
17     """Draw detected corners on image"""
18     img_with_corners = img.copy()
19     cv2.drawChessboardCorners(img_with_corners, pattern_size, corners, True)
20
21     return img_with_corners
```

## Exercise 1.3: Corner Detection Analysis

### Questions:

- What happens if the checkerboard is blurry? Test with different image qualities.
- How does corner detection accuracy affect calibration results?
- Implement a function to measure corner detection repeatability across multiple views.

# Zhang's Camera Calibration Algorithm

## Demo: Complete Calibration Pipeline

```
1 def calibrate_camera(calibration_images, pattern_size, square_size):
2     """Implement Zhang's camera calibration method"""
3
4     # Prepare object points (3D points in real world space)
5     objp = np.zeros((pattern_size[0] * pattern_size[1], 3), np.float32)
6     objp[:, :2] = np.mgrid[0:pattern_size[0], 0:pattern_size[1]].T.reshape(-1, 2)
7     objp *= square_size
8
9     # Arrays to store object points and image points
10    objpoints = [] # 3D points in real world space
11    imgpoints = [] # 2D points in image plane
12
13    for img_path in calibration_images:
14        img = load_image(img_path)
15        ret, corners = detect_checkerboard_corners(img, pattern_size)
16
17        if ret:
18            objpoints.append(objp)
19            imgpoints.append(corners)
20
21    # Perform calibration
22    ret, mtx, dist, rvecs, tvecs = cv2.calibrateCamera(
23        objpoints, imgpoints, img.shape[:2][::-1], None, None)
24
25    return ret, mtx, dist, rvecs, tvecs
```

# Calibration Analysis and Validation

## Exercise 1.4: Calibration Quality Assessment

**Task:** Implement calibration quality metrics

```
1 def compute_reprojection_error(objpoints, imgpoints, mtx, dist, rvecs, tvecs):
2     """Compute RMS reprojection error"""
3     # TODO: Project 3D points back to image plane
4     # TODO: Compute distance between projected and detected points
5     pass
6
7 def analyze_calibration_accuracy(mtx, dist, rvecs, tvecs, objpoints, imgpoints):
8     """Comprehensive calibration analysis"""
9     # TODO: Compute per-image errors
10    # TODO: Visualize error distribution
11    # TODO: Check for systematic biases
12    pass
13
```

### Questions:

- What's a "good" reprojection error? How does it depend on image resolution?
- How many calibration images do you need for stable results?
- What happens if you use only images from one orientation?

### Expected Results

- Reprojection error: < 1.0 pixels for good calibration
- Focal lengths should be approximately equal for square pixels
- Principal point should be near image center

# Homography Estimation with DLT

## Demo: Direct Linear Transform Implementation

```
1 def estimate_homography_dlt(src_pts, dst_pts):
2     """Estimate homography using Direct Linear Transform"""
3     assert len(src_pts) >= 4, "Need at least 4 point correspondences"
4
5     # Normalize points for numerical stability
6     src_norm, T1 = normalize_points(src_pts)
7     dst_norm, T2 = normalize_points(dst_pts)
8
9     # Set up linear system Ah = 0
10    A = []
11    for i in range(len(src_norm)):
12        x, y = src_norm[i]
13        u, v = dst_norm[i]
14
15        A.append([0, 0, 0, -x, -y, -1, v*x, v*y, v])
16        A.append([x, y, 1, 0, 0, 0, -u*x, -u*y, -u])
17
18    A = np.array(A)
19
20    # Solve using SVD
21    U, S, Vt = np.linalg.svd(A)
22    h = Vt[-1, :] / Vt[-1, -1] # Normalize by last element
23    H_norm = h.reshape(3, 3)
24
25    # Denormalize
26    H = np.linalg.inv(T2) @ H_norm @ T1
27
28    return H / H[2, 2] # Normalize by bottom-right element
29
```

# Point Normalization for Numerical Stability

## Demo: Hartley Normalization

```
1 def normalize_points(points):
2     """Normalize points for DLT numerical stability"""
3     points = np.array(points, dtype=np.float32)
4
5     # Compute centroid
6     centroid = np.mean(points, axis=0)
7
8     # Compute mean distance to centroid
9     distances = np.linalg.norm(points - centroid, axis=1)
10    scale = np.sqrt(2) / np.mean(distances)
11
12    # Create transformation matrix
13    T = np.array([
14        [scale, 0, -scale * centroid[0]],
15        [0, scale, -scale * centroid[1]],
16        [0, 0, 1]
17    ])
18
19    # Apply transformation
20    points_hom = np.column_stack([points, np.ones(len(points))])
21    points_norm = (T @ points_hom.T).T
22    points_norm = points_norm[:, :2] / points_norm[:, [2]]
23
24    return points_norm, T
25
```

## Exercise 1.5: Homography Robustness

**Task:** Compare normalized vs. unnormalized DLT

- Implement DLT without normalization
- Test on points with different scales (e.g., [0,1] vs [0,1000])
- Measure conditioning of matrix  $A^T A$

**Question:** Why does normalization improve numerical stability?

# RANSAC for Robust Homography Estimation

Demo: RANSAC Implementation

```
1 def ransac_homography(src_pts, dst_pts, threshold=3.0, max_iters=1000):
2     """Robust homography estimation using RANSAC"""
3     best_H = None
4     best_inliers = []
5     best_score = 0
6
7     n_points = len(src_pts)
8
9     for _ in range(max_iters):
10        # Randomly sample 4 points
11        sample_idx = np.random.choice(n_points, 4, replace=False)
12        sample_src = src_pts[sample_idx]
13        sample_dst = dst_pts[sample_idx]
14
15        # Estimate homography from sample
16        H = estimate_homography_dlt(sample_src, sample_dst)
17
18        # Count inliers
19        inliers = []
20        for i in range(n_points):
21            # Transform source point
22            src_hom = np.append(src_pts[i], 1)
23            dst_pred_hom = H @ src_hom
24            dst_pred = dst_pred_hom[:2] / dst_pred_hom[2]
25
26            # Compute reprojection error
27            error = np.linalg.norm(dst_pred - dst_pts[i])
28            if error < threshold:
29                inliers.append(i)
30
31        # Update best model if better
32        if len(inliers) > best_score:
33            best_score = len(inliers)
34            best_inliers = inliers
35            best_H = H
36
37    return best_H, best_inliers
```

# Lens Distortion Correction

Demo: Brown-Conrady Distortion Model

```
1 def undistort_points(points, mtx, dist):
2     """Undistort points using camera parameters"""
3     # Convert to homogeneous coordinates
4     points_hom = np.column_stack([points, np.ones(len(points))])
5
6     # Apply inverse camera matrix to get normalized coordinates
7     points_norm = np.linalg.inv(mtx) @ points_hom.T
8     points_norm = points_norm[:, :2].T
9
10    # Apply distortion correction
11    k1, k2, p1, p2, k3 = dist.ravel()
12
13    points_corrected = []
14    for x, y in points_norm:
15        r2 = x*x + y*y
16
17        # Radial distortion correction
18        radial_factor = 1 + k1*r2 + k2*r2*r2 + k3*r2*r2*r2
19
20        # Tangential distortion correction
21        x_corrected = x * radial_factor + 2*p1*x*y + p2*(r2 + 2*x*x)
22        y_corrected = y * radial_factor + p1*(r2 + 2*y*y) + 2*p2*x*y
23
24        points_corrected.append([x_corrected, y_corrected])
25
26    # Convert back to pixel coordinates
27    points_corrected = np.array(points_corrected)
28    points_corrected_hom = np.column_stack([points_corrected, np.ones(len(points_corrected))])
29    points_pixel = (mtx @ points_corrected_hom.T).T
30
31    return points_pixel[:, :2]
```

# Color Space Transformations

## Demo: RGB to HSV Conversion

```
1 def rgb_to_hsv_manual(rgb_img):
2     """Manual RGB to HSV conversion following lecture algorithm"""
3     rgb_normalized = rgb_img.astype(np.float32) / 255.0
4     h, w, c = rgb_img.shape
5     hsv_img = np.zeros_like(rgb_normalized)
6
7     for i in range(h):
8         for j in range(w):
9             r, g, b = rgb_normalized[i, j]
10
11             # Value (brightness)
12             v = max(r, g, b)
13             delta = v - min(r, g, b)
14
15             # Saturation
16             s = 0 if v == 0 else delta / v
17
18             # Hue
19             if delta == 0:
20                 h_val = 0
21             elif v == r:
22                 h_val = 60 * ((g - b) / delta % 6)
23             elif v == g:
24                 h_val = 60 * ((b - r) / delta + 2)
25             else: # v == b
26                 h_val = 60 * ((r - g) / delta + 4)
27
28             hsv_img[i, j] = [h_val, s, v]
29
30     return hsv_img
31
```

## Exercise 1.6: Color Space Analysis

**Task:** Compare manual vs. OpenCV color conversion

- Implement HSV to RGB conversion
- Measure numerical differences between implementations
- Visualize color space distributions for different image types

# Histogram Equalization Implementation

## Exercise 1.7: Advanced Image Enhancement

**Task:** Implement adaptive histogram equalization

```
1 def histogram_equalization(img):
2     """Global histogram equalization from scratch"""
3     # TODO: Implement algorithm from lecture
4     pass
5
6 def adaptive_histogram_equalization(img, tile_size=(8, 8)):
7     """CLAHE (Contrast Limited Adaptive Histogram Equalization)"""
8     # TODO: Implement tile-based equalization
9     # TODO: Apply contrast limiting
10    # TODO: Bilinear interpolation between tiles
11    pass
12
```

### Questions:

- When does global histogram equalization fail?
- How does tile size affect adaptive equalization results?
- What are the trade-offs between enhancement and noise amplification?

# Advanced Exercises & Extension Challenges

## Challenge Problems (Extra Credit)

For students seeking deeper understanding and implementation experience.

### Challenge 1: Multi-Camera Calibration

- Implement stereo camera calibration
- Estimate relative pose between two cameras
- Validate epipolar geometry constraints
- **Theory Connection:** How does this relate to fundamental matrix estimation?

### Challenge 2: Real-time Calibration

- Implement live camera calibration using webcam
- Add automatic checkerboard detection and quality assessment
- Visualize calibration progress in real-time
- **Theory Connection:** What's the minimum number of views needed?

### Challenge 3: Custom Calibration Patterns

- Implement calibration using circular dots instead of checkerboards
- Compare accuracy and robustness vs. checkerboard patterns
- Handle partial pattern visibility
- **Theory Connection:** How does pattern choice affect corner detection accuracy?

# Theoretical Questions for Deeper Understanding

## Conceptual Questions

- ① **Camera Model:** Why is the pinhole camera model insufficient for real cameras? When would you use fisheye vs. perspective models?
- ② **Calibration Theory:** Prove that you need at least 6 point correspondences to estimate all camera parameters. What happens with exactly 6 points?
- ③ **Homography Constraints:** Under what conditions is a homography undefined? How many degrees of freedom does a homography have and why?
- ④ **Distortion Models:** Why do lens distortion coefficients have different effects? When would higher-order terms be necessary?
- ⑤ **Numerical Stability:** Explain why the DLT algorithm becomes unstable without normalization. What is the condition number measuring?

## Discussion Questions

- How would you modify Zhang's method for a moving camera setup?
- What calibration accuracy is needed for different applications (AR, robotics, measurement)?
- How do you validate calibration quality without ground truth?

# Lab Assessment & Deliverables

## What to Submit

- ① **Implemented Functions:** All exercise functions with proper documentation
- ② **Calibration Results:** Camera parameters from your own calibration dataset
- ③ **Analysis Report:**
  - Reprojection error analysis
  - Comparison of different algorithms (DLT vs RANSAC)
  - Discussion of failure cases and limitations
- ④ **Visualization:** Clear plots showing calibration quality and algorithm behavior

## Grading Criteria

- **Correctness (40%):** Algorithms produce mathematically correct results
- **Code Quality (25%):** Clean, well-documented, efficient implementation
- **Analysis (25%):** Thoughtful discussion of results and theoretical connections
- **Creativity (10%):** Extension challenges and novel insights

## Next Lab Preview

**Lab 2:** Custom convolution operations, data augmentation, preprocessing pipelines - building the foundation for deep learning approaches.