

Lab 3: Architecture Implementation Mastery

From ResNet to Vision Transformers - Hands-On

Prof. David Olivieri

Artificial Vision (VIAR25/26) - UVigo

September 22, 2025

Lab 3 Overview: Architecture Implementation Mastery

Learning Objectives

By the end of this lab, you will be able to:

- **Implement** ResNet from mathematical principles with proper gradient flow analysis
- **Design** and integrate attention mechanisms (SE, CBAM) into existing architectures
- **Build** Vision Transformer components from scratch with position encoding
- **Compare** CNN vs ViT performance across different data regimes and computational constraints
- **Analyze** architectural trade-offs and make informed deployment decisions

Lab Structure (4 hours)

Part 1: ResNet Deep Dive (90 min)

- Mathematical foundation review
- Basic vs Bottleneck block implementation
- Gradient flow visualization
- Depth scaling experiments

Part 2: Attention Integration (90 min)

- SE module implementation
- CBAM spatial + channel attention
- Performance benchmarking
- Attention visualization

Continued...

Part 3: Vision Transformer (60 min)

- Patch embedding implementation
- Multi-head self-attention
- Complete ViT architecture

Part 4: Comparative Analysis (60 min)

- CNN vs ViT benchmarking
- Computational profiling
- Architecture selection framework

Tools & Environment

- PyTorch 2.0+, torchvision
- CIFAR-10/100 datasets
- GPU recommended (Google Colab OK)
- Weights & Biases for experiment tracking

Section 1

Part 1: ResNet Deep Dive

Part 1: ResNet Implementation Deep Dive

1.1 Mathematical Foundation Review (15 min)

- Derive residual learning formulation: $H(x) = F(x) + x$
- Analyze gradient flow: $\frac{\partial \mathcal{L}}{\partial x} = \frac{\partial \mathcal{L}}{\partial H} \cdot (1 + \frac{\partial F}{\partial x})$
- Identify mapping importance and degradation problem

1.2 Block Implementation (45 min)

- BasicBlock: $3 \times 3 \rightarrow 3 \times 3$ convolutions with skip connection
- BottleneckBlock: $1 \times 1 \rightarrow 3 \times 3 \rightarrow 1 \times 1$ with $4 \times$ channel reduction
- Pre-activation vs post-activation comparison
- Proper padding and stride handling for dimension matching

1.3 Architecture Scaling (30 min)

- Build ResNet-18, 34, 50, 101 from block components
- Measure training time, memory usage, and convergence
- Plot accuracy vs depth relationship
- Analyze when deeper networks help vs hurt

ResNet Implementation: Key Components

Basic Block Structure

```

1 class BasicBlock(nn.Module):
2     def __init__(self, in_channels, out_channels,
3                  stride=1, downsample=None):
4         # TODO: Implement basic block
5         # - Two 3x3 convolutions
6         # - Batch normalization
7         # - ReLU activation
8         # - Skip connection handling
9         pass
10
11    def forward(self, x):
12        # TODO: Implement forward pass
13        # Remember: out = F(x) + x
14        pass
15

```

Bottleneck Block Structure

```

1 class BottleneckBlock(nn.Module):
2     expansion = 4
3
4     def __init__(self, in_channels, out_channels,
5                  stride=1, downsample=None):
6         # TODO: Implement bottleneck block
7         # - 1x1 conv (reduce channels)
8         # - 3x3 conv (spatial processing)
9         # - 1x1 conv (expand channels)
10        # - Batch norm + ReLU between each
11        pass
12
13    def forward(self, x):
14        # TODO: Implement forward pass
15        pass
16

```

Key Implementation Points

- Handle dimension mismatch with 1×1 convolution in skip connection
- Proper stride application only on first convolution of each block
- BatchNorm placement: after conv, before ReLU (post-activation ResNet)
- Expansion factor of 4 for bottleneck blocks

Section 2

Part 2: Attention Integration

Part 2: Attention Mechanism Integration

2.1 SE Module Implementation (30 min)

- Global average pooling: $z_c = \frac{1}{H \times W} \sum_{i,j} x_{i,j,c}$
- Excitation network: $s = \sigma(W_2 \delta(W_1 z))$ with reduction ratio
- Channel-wise multiplication: $\tilde{x}_{i,j,c} = s_c \cdot x_{i,j,c}$
- Integration into ResNet blocks

2.2 CBAM Implementation (30 min)

- Channel attention: Apply SE mechanism
- Spatial attention: $M_s = \sigma(\text{conv77}([\text{AvgPool}(F); \text{MaxPool}(F)]))$
- Sequential application: $F'' = M_s(F') \odot F'$ where $F' = M_c(F) \odot F$

2.3 Performance Analysis (30 min)

- Benchmark ResNet vs ResNet+SE vs ResNet+CBAM
- Measure computational overhead and memory usage
- Visualize attention maps using Grad-CAM
- Analyze which channels/regions receive highest attention

Attention Mechanisms: SE and CBAM

SE Module

```

1 class SEModule(nn.Module):
2     def __init__(self, channels, reduction=16):
3         super().__init__()
4         # TODO: Implement SE module
5         # - Global Average Pooling
6         # - FC layer with reduction
7         # - ReLU activation
8         # - FC layer back to channels
9         # - Sigmoid activation
10        pass
11
12    def forward(self, x):
13        # TODO: Implement forward pass
14        # 1. Global avg pool (B,C,H,W) -> (B,C,1,1)
15        # 2. Squeeze through FC layers
16        # 3. Sigmoid to get weights
17        # 4. Scale input: x * weights
18        pass
19

```

CBAM Module

```

1 class CBAM(nn.Module):
2     def __init__(self, channels, reduction=16):
3         super().__init__()
4         # TODO: Implement CBAM
5         # - Channel attention (SE-like)
6         # - Spatial attention
7         # * Avg + Max pool along channel
8         # * 7x7 conv + sigmoid
9         pass
10
11    def forward(self, x):
12        # TODO: Implement forward pass
13        # 1. Apply channel attention
14        # 2. Apply spatial attention
15        # 3. Return attended features
16        pass
17

```

Implementation Tips

- Use AdaptiveAvgPool2d(1) for global pooling
- Reduction ratio typically 16 for SE modules
- For CBAM spatial attention, concatenate avg and max pooled features
- Visualize attention maps with torch.nn.functional.interpolate

Section 3

Part 3: Vision Transformer

Part 3: Vision Transformer from Scratch

3.1 Patch Embedding Layer (20 min)

- Image to patches: $(H, W, C) \rightarrow (N, P^2 \cdot C)$ where $N = \frac{H \times W}{P^2}$
- Linear projection: $\mathbf{x}_i \rightarrow \mathbf{e}_i = \mathbf{x}_i \mathbf{W}_e + \mathbf{b}_e$
- Class token prepending: $[\mathbf{x}_{cls}, \mathbf{e}_1, \dots, \mathbf{e}_N]$
- Position encoding addition (learnable vs sinusoidal)

3.2 Multi-Head Self-Attention (25 min)

- Query, Key, Value projections for each head
- Scaled dot-product attention: $\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$
- Multi-head concatenation and output projection
- Attention weight visualization and interpretation

3.3 Complete ViT Architecture (15 min)

- Transformer encoder block: LayerNorm \rightarrow MSA \rightarrow Add \rightarrow LayerNorm \rightarrow MLP \rightarrow Add
- Stack multiple encoder layers
- Classification head: Extract [CLS] token \rightarrow Linear layer
- Hyperparameter exploration: patch size, embedding dimension, number of heads

Vision Transformer: Core Components

Patch Embedding

```

1 class PatchEmbedding(nn.Module):
2     def __init__(self, img_size=224, patch_size=16,
3                  in_channels=3, embed_dim=768):
4         super().__init__()
5         # TODO: Implement patch embedding
6         # — Calculate number of patches
7         # — Conv2d with kernel=patch_size, stride=patch_size
8         # — Or use nn.Unfold + Linear
9         pass
10
11    def forward(self, x):
12        # TODO: Implement forward pass
13        # (B, C, H, W) -> (B, N, embed_dim)
14        # where N = (H*W)/(patch_size^2)
15        pass
16

```

Multi-Head Attention

```

1 class MultiHeadAttention(nn.Module):
2     def __init__(self, embed_dim, num_heads):
3         super().__init__()
4         # TODO: Implement multi-head attention
5         # — Q, K, V linear projections
6         # — Output projection
7         # — Proper head dimension calculation
8         pass
9
10
11    def forward(self, x):
12        # TODO: Implement forward pass
13        # 1. Linear projections to Q, K, V
14        # 2. Reshape for multi-head
15        # 3. Scaled dot-product attention
16        # 4. Concatenate heads
17        # 5. Output projection
18        pass
19
20

```

ViT Implementation Details

- Use `nn.Conv2d(kernel_size=patch_size, stride=patch_size)` for patch embedding
- Add learnable CLS token and position embeddings
- Scale attention by $\sqrt{d_k}$ for stability
- Use GELU activation in MLP blocks
- Apply LayerNorm before attention and MLP (pre-norm)

Section 4

Part 4: Comparative Analysis

Part 4: Comprehensive Comparative Analysis

4.1 Performance Benchmarking (25 min)

- Train ResNet-50 and ViT-Base on CIFAR-10/100
- Measure training convergence, final accuracy, inference speed
- Vary dataset size: 1K, 10K, full dataset
- Plot learning curves and analyze data efficiency

4.2 Computational Profiling (20 min)

- Memory usage: Peak GPU memory during training/inference
- FLOPs analysis: Theoretical and measured computational cost
- Latency benchmarking: Batch size 1 vs 32 vs 128
- Energy consumption measurement (if hardware supports)

4.3 Architecture Selection Framework (15 min)

- Create decision tree: Dataset size → Computational budget → Task requirements
- Practical deployment scenarios and recommendations
- Trade-off analysis: Accuracy vs Speed vs Memory vs Data efficiency

Benchmarking and Profiling Framework

Performance Measurement

```
1 class ModelProfiler:
2     def __init__(self, model, device):
3         self.model = model
4         self.device = device
5
6     def measure_flops(self, input_size):
7         # TODO: Implement FLOPs counting
8         # Use torchprofile or similar
9         pass
10
11    def measure_memory(self, batch_size, input_size):
12        # TODO: Implement memory profiling
13        # Track GPU memory usage
14        pass
15
16    def measure_latency(self, input_tensor, num_runs=100):
17        # TODO: Implement latency measurement
18        # Average over multiple runs
19        pass
20
21    def benchmark_training(self, dataloader, epochs=5):
22        # TODO: Implement training benchmark
23        # Track convergence speed
24        pass
25
```

Comparison Metrics

- **Accuracy:** Top-1 and Top-5 on test set
- **Training time:** Time to reach target accuracy
- **Memory:** Peak GPU memory usage
- **FLOPs:** Computational complexity
- **Latency:** Inference time per image
- **Data efficiency:** Performance vs dataset size

Lab Assessment and Deliverables

Coding Assessment (60%)

Implementation Quality:

- Correct mathematical implementation
- Clean, well-documented code
- Proper error handling and edge cases
- Efficient memory and computational usage

Architecture Components:

- ResNet blocks (basic & bottleneck)
- SE and CBAM attention modules
- ViT patch embedding and self-attention
- Proper integration and testing

Analysis Report (30%)

Required Sections:

- ① ResNet depth scaling analysis
- ② Attention mechanism effectiveness
- ③ CNN vs ViT performance comparison
- ④ Computational trade-off analysis
- ⑤ Architecture selection recommendations

Presentation Demo (10%)

5-minute presentation including:

- Live demonstration of implemented architectures
- Key findings from comparative analysis
- Most surprising/interesting result
- Practical deployment recommendation

Submission Details

Due Date: 1 week from lab session

Format:

- Single Jupyter notebook with all code
- Embedded markdown analysis sections
- All plots and visualizations included
- Runnable on standard GPU hardware

Bonus Points (5%): Creative extensions or optimizations beyond requirements

Getting Started

- Download template from course website
- Set up environment: `pip install -r requirements.txt`
- Start with Part 1: ResNet implementation
- Use provided test cases to verify implementations