

Artificial Vision (VIAR25/26)

Lab 5: Advanced Image Segmentation

Prof. David Olivieri

UVigo

October 2, 2025

Lab Objectives

Today's Mission

Implement semantic segmentation architectures from scratch, master advanced loss functions, and build a trainable Mini-SAM model for prompt-based segmentation.

Learning Outcomes

- ① **Build** FCN-8s from pretrained ResNet with skip connections
- ② **Implement** DeepLabV3+ with ASPP module and atrous convolutions
- ③ **Master** advanced loss functions: Dice, Focal, combined strategies
- ④ **Train** Mini-SAM from scratch with point/box prompting
- ⑤ **Analyze** architectural trade-offs and performance metrics

Datasets

- **PASCAL VOC 2012:** 21 classes, semantic segmentation benchmark
- **VOC Subset:** 500 images for Mini-SAM training
- **Synthetic shapes:** Quick testing and debugging



Lab Structure

Part 1: FCN Architecture (40 min)

- **Task 1.1:** Convert ResNet to fully convolutional network
- **Task 1.2:** Implement score layers and progressive upsampling
- **Task 1.3:** Build FCN-32s, FCN-16s, FCN-8s with skip connections
- **Task 1.4:** Train and compare skip connection impact

Part 2: DeepLabV3+ with ASPP (40 min)

- **Task 2.1:** Implement atrous convolution operations
- **Task 2.2:** Build complete ASPP module with 5 branches
- **Task 2.3:** Create encoder-decoder with low-level feature fusion
- **Task 2.4:** Compare performance with FCN variants

Part 3: Mini-SAM Training from Scratch (40 min)

- **Task 3.1:** Build lightweight Mini-SAM architecture
- **Task 3.2:** Implement prompt encoding (points and boxes)
- **Task 3.3:** Train Mini-SAM with simulated prompts
- **Task 3.4:** Comparative analysis of all methods

Part 1: Fully Convolutional Networks (FCN)

From Classification to Dense Prediction

Transform a pretrained ResNet into a segmentation network by replacing fully connected layers with 1×1 convolutions and adding skip connections to recover spatial detail.

FCN Architecture Variants

- **FCN-32s:** Single $32 \times$ upsampling from conv5
- **FCN-16s:** Add skip from conv4 (stride 16)
- **FCN-8s:** Add skip from conv3 (stride 8)
- Progressive improvement in boundary detail

Implementation Components

- ① Load pretrained ResNet50
- ② Extract features at multiple scales
- ③ Add 1×1 score layers for classification
- ④ Implement transposed convolutions for upsampling
- ⑤ Fuse features with element-wise addition

Key Mathematical Insight

FC layer as convolution:

$$\mathbf{W}_{\text{FC}} \in \mathbb{R}^{K \times D} \equiv \mathbf{W}_{\text{conv}} \in \mathbb{R}^{K \times H' \times W' \times C}$$

Enables arbitrary input sizes!

Expected Performance (PASCAL VOC)

- FCN-32s: 59% mIoU
- FCN-16s: 61% mIoU
- FCN-8s: 63% mIoU
- Each skip adds 2% mIoU

Part 1: FCN Code Structure

```
1 class FCN8s(nn.Module):
2     """TODO: Implement FCN-8s with skip connections"""
3     def __init__(self, n_classes=21):
4         super().__init__()
5         # Task 1.1: Load pretrained ResNet50
6         resnet = models.resnet50(pretrained=True)
7
8         # Extract encoder layers
9         self.conv1 = resnet.conv1
10        self.bn1 = resnet.bn1
11        self.relu = resnet.relu
12        self.maxpool = resnet.maxpool
13        self.layer1 = resnet.layer1 # 1/4
14        self.layer2 = resnet.layer2 # 1/8 - pool3 for skip
15        self.layer3 = resnet.layer3 # 1/16 - pool4 for skip
16        self.layer4 = resnet.layer4 # 1/32
17
18        # Task 1.2: Add score layers (1x1 conv to n_classes)
19        self.score_pool3 = None # TODO: Conv2d(1024, n_classes, 1)
20        self.score_pool4 = None # TODO: Conv2d(2048, n_classes, 1)
21        self.score_fr = None # TODO: Conv2d(2048, n_classes, 1)
22
23        # Task 1.3: Add upsampling layers
24        self.upscore2 = None # TODO: ConvTranspose2d
25        self.upscore_pool4 = None # TODO: ConvTranspose2d
26        self.upscore8 = None # TODO: ConvTranspose2d
27
28    def forward(self, x):
29        # Task 1.4: Implement forward pass with progressive skip fusion
30        pass
```

Part 1: FCN-8s Forward Pass — Pseudocode + Intuition

Algorithm 1: FCN-8s Forward Pass with Skip Connections

```
Data: Input image I ∈ ℝB×3×H×W
Result: Segmentation map S ∈ ℝB×K×H×W
Encoder Path:
x ← relu(bn1(conv1(I)));
x ← maxpool(x);
x ← layer1(x) // stride 4
pool3 ← layer2(x) // stride 8, save for skip
pool4 ← layer3(pool3) // stride 16, save for skip
x ← layer4(pool4) // stride 32
Score Layers (1x1 convolutions):
score_fr ← Conv1x1(x, 2048 → K);
score_pool4 ← Conv1x1(pool4, 1024/2048 → K);
score_pool3 ← Conv1x1(pool3, 512/1024 → K);
Progressive Upsampling with Skips:
// First skip: pool4 at stride 16
upscore2 ← ConvTranspose2d(score_fr, stride = 2) // 32 → 16
fuse_pool4 ← upscore2 + score_pool4 // elementwise add
// Second skip: pool3 at stride 8
upscore_pool4 ← ConvTranspose2d(fuse_pool4, stride = 2) // 16 → 8
fuse_pool3 ← upscore_pool4 + score_pool3 // elementwise add
// Final upsampling to original resolution
S ← ConvTranspose2d(fuse_pool3, stride = 8) // 8 → 1 (H,W)
return S;
```

PyTorch Implementation Notes

- **1x1 score layers:** nn.Conv2d(C_in, K, kernel_size=1) maps feature channels → classes.
- **Upsampling:** Use nn.ConvTranspose2d(K,K,4,2,1) for ×2; for ×8: kernel_size=16,stride=8,padding=4.
- **Shape alignment:** Ensure skip tensors have same H×W and channels = K before addition.
- **Alt. upsampling:**
F.interpolate(..., scale_factor=2, mode='bilinear', align_corners=False) + 1×1 convs is OK (simpler, fewer params).
- **Padding/stride:** Mismatched sizes? Use center_crop or F.pad before addition.
- **Normalization:** Keep the ResNet stem (conv1/bn1/relu/maxpool) as-is for correct strides.
- **Initialization:** Kaiming for new convs; optionally freeze early ResNet layers then unfreeze.

Typical Tensor Shapes (example H=W=512)

I : B × 3 × 512 × 512
layer1 : B × 256 × 128 × 128 (/4)
pool3 = layer2 : B × 512 × 64 × 64 (/8)
pool4 = layer3 : B × 1024 × 32 × 32 (/16)
layer4 : B × 2048 × 16 × 16 (/32)
S : B × K × 512 × 512

Common Pitfalls

- Adding tensors with different H×W (off-by-one from padding).
- Forgetting to set bias=False on 1x1 score layers if followed by BN (optional).
- Training only with CrossEntropy → slow boundary recovery; try CE + Dice.
- Class imbalance in VOC: consider ignore_index=255 and class weights.

Part 2: DeepLabV3+ with ASPP

Multi-Scale Context via Atrous Convolutions

ASPP (Atrous Spatial Pyramid Pooling) captures multi-scale context without losing resolution by using parallel atrous convolutions at different dilation rates.

Atrous Convolution

Standard conv:

$$y[i, j] = \sum_{m,n} x[i + m, j + n] \cdot w[m, n]$$

Atrous conv (rate r):

$$y[i, j] = \sum_{m,n} x[i + r \cdot m, j + r \cdot n] \cdot w[m, n]$$

Effective kernel size:

$$k_{\text{eff}} = k + (k - 1)(r - 1)$$

For 3×3 kernel: $r = 6 \rightarrow 13 \times 13$ effective

ASPP Module (5 branches)

- ❶ 1×1 convolution (point-wise features)
- ❷ 3×3 atrous conv, rate=6
- ❸ 3×3 atrous conv, rate=12
- ❹ 3×3 atrous conv, rate=18
- ❺ Global average pooling + 1×1 conv

All outputs concatenated $\rightarrow 1 \times 1$ projection

DeepLabV3+ Decoder

- $4 \times$ bilinear upsample from ASPP
- Concatenate with low-level features (stride 4)
- 3×3 conv refinement
- Final $4 \times$ upsampling to original size

Part 2: ASPP Module Code Structure

```
1 class ASPP(nn.Module):
2     """TODO: Implement Atrous Spatial Pyramid Pooling"""
3     def __init__(self, in_channels, out_channels=256, rates=[6, 12, 18]):
4         super().__init__()
5
6         # Task 2.1: Branch 1 - 1x1 convolution
7         self.conv1 = None # TODO: Sequential(Conv2d, BatchNorm, ReLU)
8
9         # Task 2.1: Branches 2-4 - Atrous convolutions
10        self.atrous.convs = nn.ModuleList()
11        # TODO: for rate in rates: append Conv2d with dilation=rate
12
13        # Task 2.1: Branch 5 - Global pooling
14        self.global.avg.pool = None # TODO: AdaptiveAvgPool + Conv
15
16        # Task 2.2: Fusion layer (5 * out_channels -> out_channels)
17        self.conv.out = None # TODO: Conv + BN + ReLU + Dropout
18
19    def forward(self, x):
20        # Task 2.2: Apply all branches and concatenate
21        pass
```

Part 2: ASPP — Pseudocode + Intuition

Algorithm 2: ASPP Module Implementation

```
Data: Feature map  $\mathbf{X} \in \mathbb{R}^{B \times C \times H \times W}$ , Rates {6, 12, 18}
Result: Multi-scale features  $\mathbf{Y} \in \mathbb{R}^{B \times 256 \times H \times W}$ 
size  $\leftarrow (H, W)$  // Save spatial dimensions
Branch 1:  $1 \times 1$  Convolution
 $\mathbf{f}_0 \leftarrow \text{Conv}_{1 \times 1}(\mathbf{X}, C \rightarrow 256);$ 
 $\mathbf{f}_0 \leftarrow \text{BN}(\mathbf{f}_0) \rightarrow \text{ReLU}(\mathbf{f}_0);$ 
Branches 2–4: Atrous Convolutions
for  $i = 1$  to 3 do
     $r_i \leftarrow \{6, 12, 18\}[i]$  // Dilation rate
     $\mathbf{f}_i \leftarrow \text{Conv}_{3 \times 3}(\mathbf{X}, \text{dilation} = r_i, \text{padding} = r_i);$ 
     $\mathbf{f}_i \leftarrow \text{BN}(\mathbf{f}_i) \rightarrow \text{ReLU}(\mathbf{f}_i);$ 
end
Branch 5: Global Average Pooling
 $\mathbf{f}_g \leftarrow \text{AdaptiveAvgPool2d}(\mathbf{X}, 1);$ 
 $\mathbf{f}_g \leftarrow \text{Conv}_{1 \times 1}(\mathbf{f}_g, C \rightarrow 256);$ 
 $\mathbf{f}_g \leftarrow \text{BN}(\mathbf{f}_g) \rightarrow \text{ReLU}(\mathbf{f}_g);$ 
 $\mathbf{f}_4 \leftarrow \text{Interpolate}(\mathbf{f}_g, \text{size} = (H, W));$ 
Feature Fusion
 $\mathbf{Y}_{\text{cat}} \leftarrow \text{Concat}([\mathbf{f}_0, \mathbf{f}_1, \mathbf{f}_2, \mathbf{f}_3, \mathbf{f}_4], \text{dim} = 1);$ 
 $\mathbf{Y} \leftarrow \text{Conv}_{1 \times 1}(\mathbf{Y}_{\text{cat}}, 1280 \rightarrow 256);$ 
 $\mathbf{Y} \leftarrow \text{BN}(\mathbf{Y}) \rightarrow \text{ReLU}(\mathbf{Y}) \rightarrow \text{Dropout}(\mathbf{Y}, p = 0.1);$ 
return  $\mathbf{Y}$ 
```

PyTorch Implementation Notes

- **Dilated 3×3 :** set padding=dilation to keep $H \times W$.
- **Rates vs output stride:** with $OS=16$, typical rates are {6,12,18}; with $OS=8$, halve them (e.g., {3,6,9}) to avoid gridding.
- **Branch 5 upsampling:**
 $F.interpolate(\mathbf{f}_g, \text{size}=(H,W), \text{mode}=\text{'bilinear'}, \text{align_corners=False})$.
- **Concat dim:** channel axis = dim=1 $\rightarrow 5 \times 256 = 1280$ before the 1×1 projection.
- **Separable conv (optional):** DeepLabV3+ often uses depthwise-separable 3×3 for speed/memory.
- **Norm choice:** small batch? prefer SyncBN/GroupNorm over BN to stabilize training.
- **Dropout:** light regularization ($p=0.1$) after the projection is standard.

Typical Shapes (ResNet, OS=16)

$\mathbf{X} : B \times C \times 32 \times 32$ ($H=W=32$ if input 512)
 $\mathbf{f}_0 \dots \mathbf{f}_4 : B \times 256 \times 32 \times 32$
 $\mathbf{Y}_{\text{cat}} : B \times 1280 \times 32 \times 32$
 $\mathbf{Y} : B \times 256 \times 32 \times 32$

Common Pitfalls

- **Wrong padding** with dilation \Rightarrow spatial mismatch on concat.
- **Too-large rates** at small $H, W \Rightarrow$ aliasing/gridding; adjust by OS.
- **align_corners=True** can distort scaling; keep **False** unless you know why.
- **Forgetting to project low-level features** (e.g., 48 ch) in the decoder before concat.

Part 3: Mini-SAM - Trainable from Scratch

Educational Prompt-Based Segmentation

Mini-SAM is a simplified, trainable version of SAM with 5M parameters (vs 636M) that can be trained from scratch on limited data while demonstrating the key concepts of prompt-based segmentation.

Mini-SAM Architecture

Image Encoder (Lightweight):

- MobileNetV3-Small backbone
- Output: $H/8 \times W/8 \times 256$ feature map
- Fast encoding (50ms)

Prompt Encoder:

- Point encoding: type + position embeddings
- Box encoding: corner coordinates
- Broadcast to spatial dimensions

Mask Decoder (Conv-based):

- Fuse image + prompt features
- 3-layer CNN decoder
- Output: class logits + IoU prediction

Training Strategy

Simulated Prompts:

- Sample points from ground truth masks
- 50% foreground, 50% background points
- Random number of points (3-10)
- Optional box prompts from mask bounds

Loss Function:

$$\mathcal{L} = \mathcal{L}_{CE} + \mathcal{L}_{Dice} + 0.1 \cdot \mathcal{L}_{IoU}$$

Training Data:

- PASCAL VOC subset: 500 images
- Augmentation: flip, scale, crop
- Training time: 2-3 hours on single GPU

Key Advantages

Part 3: Mini-SAM Code Structure

```
1 class MiniSAM(nn.Module):
2     """TODO: Implement lightweight trainable SAM"""
3     def __init__(self, n_classes=21, embed_dim=256):
4         super().__init__()
5
6         # Task 3.1: Lightweight Image Encoder (~2M params)
7         backbone = models.mobilenet_v3_small(pretrained=True)
8         self.image_encoder = None # TODO: Extract features
9         self.img_proj = None      # TODO: Project to embed_dim
10
11        # Task 3.2: Simple Prompt Encoder
12        self.point_type_embed = None # TODO: Embedding(2, embed_dim)
13        self.point_pos_embed = None # TODO: MLP for position
14        self.box_embed = None       # TODO: MLP for boxes
15
16        # Task 3.3: Decoder
17        self.decoder = None        # TODO: Conv layers for fusion
18        self.mask_head = None      # TODO: Output n_classes
19        self.iou_head = None       # TODO: IoU prediction
20
21    def forward(self, images, points=None, point_labels=None, boxes=None):
22        # Task 3.4: Implement forward pass
23        # 1. Encode image
24        # 2. Encode prompts
25        # 3. Fuse features
26        # 4. Decode mask + IoU
27        pass
```

Part 3: Mini-SAM — Training Loop & Intuition

Algorithm 3: Mini-SAM Training with Simulated Prompts

```
Data: Images I, Ground truth masks  $M_{gt}$ 
Result: Trained Mini-SAM model
Model Initialization:
Initialize MobileNetV3 encoder (pretrained);
Initialize prompt encoders (random);
Initialize decoder (random);

Training Loop:
for each batch  $(I_b, M_{gt,b})$  do
    // Simulate user prompts from ground truth
     $P, L \leftarrow \text{SamplePoints}(M_{gt,b}, n \in [3, 10], p_{fg}=0.5);$ 
     $B \leftarrow \text{SampleBoxesFromMasks}(M_{gt,b})$  // optional

    // Forward pass
     $F_{img} \leftarrow \text{ImageEncoder}(I_b);$ 
     $F_{prompt} \leftarrow \text{PromptEncoder}(P, L, B);$ 
     $F_{fused} \leftarrow \text{Concat}(F_{img}, F_{prompt});$ 
     $(M_{pred}, IoU) \leftarrow \text{Decoder}(F_{fused});$ 

    // Compute losses
     $\mathcal{L}_{CE} \leftarrow \text{CrossEntropy}(M_{pred}, M_{gt,b});$ 
     $\mathcal{L}_{Dice} \leftarrow 1 - \text{DiceCoeff}(M_{pred}, M_{gt,b});$ 
     $\text{IoU}_{true} \leftarrow \text{ComputeIoU}(M_{pred}, M_{gt,b});$ 
     $\mathcal{L}_{IoU} \leftarrow \text{MSE}(\text{IoU}, \text{stopgrad}(\text{IoU}_{true}));$ 
     $\mathcal{L}_{total} \leftarrow \mathcal{L}_{CE} + \mathcal{L}_{Dice} + 0.1 \cdot \mathcal{L}_{IoU};$ 

    // Backward pass and update
    optimizer.zero_grad();
     $\mathcal{L}_{total}.backward();$ 
    clip_grad_norm_( $\theta, \tau$ ) // e.g.,  $\tau = 1.0$ 
    optimizer.step();
    scheduler.step( $\mathcal{L}_{val}$ ) // optional
end

return Trained model;
```

PyTorch Implementation Notes

- **Batching prompts:** fix N_{pts} per image (pad/truncate); tensors points: $(B, N, 2)$ in normalized coords $[0, 1]$ or pixel (x, y) .
- **Point types:** point_labels: (B, N) with $\{0:\text{bg}, 1:\text{fg}\}$; embed via `nn.Embedding(2, D)`.
- **Box prompts:** boxes: $(B, 4)$ as $(x_{min}, y_{min}, x_{max}, y_{max})$, normalize to $[0, 1]$, project with MLP to D .
- **Broadcasting:** tile prompt embeddings to $(B, D, H/8, W/8)$ and concat with image features along dim=1.
- **Shapes (512x512, OS=8):**
 $F_{img} : B \times 256 \times 64 \times 64$, $F_{prompt} : B \times 256 \times 64 \times 64$, $F_{fused} : B \times 512 \times 64 \times 64$.
- **Losses:** `CrossEntropyLoss(ignore_index=255)` on `logits` (no softmax first). Dice on softmax probs.
- **IoU head:** predicts scalar $(B, 1)$ or per-class (B, K) ; compute `target IoU` with `no_grad`.

Stability, Speed, and Tricks

- **Warmup:** freeze backbone for 1–3 epochs, then unfreeze.
- **Optim:** AdamW ($lr=3e-4$, $wd=1e-4$); cosine or ReduceLROnPlateau.
- **AMP:** use `torch.cuda.amp.autocast` + `GradScaler` for 1.5–2x speed.
- **Grad clip:** `clip_grad_norm_` to avoid exploding grads in decoder.
- **Augment:** random flip/scale/crop; keep points/boxes consistent with transforms.

Common Pitfalls

- Mismatch between transformed images and untransformed prompts \Rightarrow wrong supervision.
- Applying softmax before `CrossEntropyLoss` (double-softmax bug).
- Computing IoU with logits instead of thresholded/argmax masks for the target signal.
- Forgetting to normalize point/box coordinates to the encoder stride (OS).

Loss Functions and Evaluation Metrics

Loss Function Implementations

Cross-Entropy: $\mathcal{L}_{CE} = -\frac{1}{N} \sum_{i,c} y_{i,c} \log(\hat{y}_{i,c})$

Dice Loss: $\mathcal{L}_{Dice} = 1 - \frac{2 \sum y \hat{y} + \epsilon}{\sum y + \sum \hat{y} + \epsilon}$

Focal Loss: $\mathcal{L}_{Focal} = -\alpha(1 - p_t)^\gamma \log(p_t)$ where $p_t = \begin{cases} p & \text{if } y = 1 \\ 1 - p & \text{else} \end{cases}$

IoU Prediction Loss: $\mathcal{L}_{IoU} = \text{MSE}(\text{IoU}_{pred}, \text{IoU}_{true})$

Combined: $\mathcal{L} = \lambda_1 \mathcal{L}_{CE} + \lambda_2 \mathcal{L}_{Dice} + \lambda_3 \mathcal{L}_{IoU}$

Evaluation Metrics

Mean IoU: $mIoU = \frac{1}{K} \sum_{k=1}^K \frac{|A_k \cap B_k|}{|A_k \cup B_k|}$

Pixel Accuracy: $PA = \frac{\sum_k n_{kk}}{\sum_k t_k}$ where n_{kk} = correctly classified pixels

IoU Accuracy: Correlation between predicted and true IoU scores

Training Configuration and Best Practices

Hyperparameters

FCN/DeepLab:

- AdamW, LR: 1e-4, weight decay 1e-4
- Batch size: 8-16
- Epochs: 50

Mini-SAM:

- AdamW, LR: 1e-4
- Batch size: 16
- Epochs: 30
- Points per sample: 3-10 (random)

Learning Rate Schedule:

- ReduceLROnPlateau with patience=5
- Or polynomial decay

Data Augmentation

- Random crop to 256×256
- Random horizontal flip ($p=0.5$)
- Random scale (0.8 to 1.2)
- Color jitter
- Normalize: ImageNet statistics

Validation Strategy

- Evaluate every epoch
- Save best mIoU checkpoint
- Track per-class IoU
- Visualize predictions
- For Mini-SAM: test with different prompt counts

Debugging Tips

- Start with 1 image overfitting
- Check tensor shapes
- Visualize prompt sampling
- Use synthetic data first

Expected Results and Performance Comparison

Quantitative Results (PASCAL VOC 2012 val / subset)

Model	mIoU (%)	Params (M)	Inference (ms)	Training Time
FCN-32s	59.4	140	35	2-3 hours
FCN-16s	61.2	140	38	2-3 hours
FCN-8s	63.1	140	42	3-4 hours
DeepLabV3+	77.5	60	80	6-8 hours
Mini-SAM (ours)	65-68	5	60	2-3 hours

For reference (not implemented in lab):

SAM (zero-shot)	75.0	636	1050	N/A
SAM (fine-tuned)	82.0	636	1050	2-3 hours

Key Observations

- **FCN:** Skip connections provide consistent improvement (+3.7% from FCN-32s to FCN-8s)
- **DeepLabV3+:** ASPP multi-scale context gives major boost (+14.4% over FCN-8s)
- **Mini-SAM:** Competitive with FCN-8s despite being trained from scratch with limited data!
- **Efficiency:** Mini-SAM has 28× fewer parameters than FCN, 127× fewer than full SAM
- **Pedagogical value:** Only Mini-SAM shows complete training pipeline

Comparative Analysis and Use Cases

When to Use Each Method

FCN-8s:

- Real-time applications (40ms inference)
- Good baseline for research
- Simple architecture for learning

DeepLabV3+:

- Production segmentation systems
- Best performance/speed trade-off
- Moderate computational budget

Mini-SAM:

- Understanding prompt-based segmentation
- Limited training data scenarios
- Interactive applications (point/box prompts)
- Rapid prototyping and experimentation
- Educational demonstrations

Mini-SAM vs Full SAM

Advantages of Mini-SAM for learning:

- Can be trained from scratch in lab timeframe
- Understandable architecture (5M params)
- Shows complete workflow: data → training → inference
- Demonstrates key concepts without requiring massive compute

Deliverables and Grading Rubric

Required Implementations (100 points)

1. FCN Implementation (25 points):

- Working FCN-32s, FCN-16s, FCN-8s (12 points)
- Training curves showing improvement (5 points)
- Qualitative visualization (4 points)
- Ablation study (4 points)

2. DeepLabV3+ Implementation (25 points):

- Complete ASPP module (12 points)
- Encoder-decoder architecture (8 points)
- Performance comparison with FCN (5 points)

3. Mini-SAM Implementation (30 points):

- Architecture implementation (10 points)
- Prompt sampling and encoding (8 points)
- Training from scratch (8 points)
- Interactive demo with point/box prompts (4 points)

4. Analysis and Report (20 points):

- Quantitative comparison table (8 points)
- Discussion of trade-offs (6 points)
- Code quality and documentation (6 points)

Submission Guidelines

What to Submit

1. Code Files:

- lab05_student.py - All implementations
- train_fcn.py - FCN training script (optional separate file)
- train_deeplab.py - DeepLab training script (optional)
- train_minisam.py - Mini-SAM training script (optional)
- demo_minisam.ipynb - Interactive demo notebook

2. Model Checkpoints:

- Best FCN-8s model (.pth)
- Best DeepLabV3+ model (.pth)
- Best Mini-SAM model (.pth)
- Training logs (TensorBoard or CSV)

3. Report (PDF):

- 3-5 pages including figures and tables
- Architecture descriptions
- Quantitative results table (all 3 methods)
- Qualitative predictions visualization
- Discussion: Why does Mini-SAM work despite limited training?

Deadline and Format

Due: Two weeks from today, 23:59

Format: Single ZIP file uploaded to course platform

Naming: lab05_[lastname]_[firstname].zip

Common Pitfalls and Solutions

FCN Issues

Dimension mismatch in skip addition:

- Check spatial sizes match exactly
- Use center cropping if needed

Checkerboard artifacts:

- Initialize transposed conv with bilinear
- Try bilinear upsample + conv instead

DeepLabV3+ Issues

ASPP checkerboard:

- Ensure padding = dilation rate
- Check all branches have BN

Out of memory:

- Reduce batch size to 4
- Lower ASPP channels to 128

Mini-SAM Issues

Poor convergence:

- Check prompt sampling (balanced fg/bg)
- Verify normalization of point coordinates
- Start with more points (8-10)
- Use pretrained MobileNet encoder

IoU prediction always wrong:

- Ensure IoU computed correctly
- Check loss weight (should be 0.1)
- Normalize IoU to [0, 1]

Features not aligning:

- Verify prompt broadcasting to spatial size
- Check concatenation dimensions
- Print intermediate shapes

Debugging strategy:

- Test with 1 class first
- Visualize sampled prompts
- Use synthetic circles/squares

Resources and References

Key Papers

- **FCN**: Long et al., "Fully Convolutional Networks for Semantic Segmentation," CVPR 2015
- **DeepLabV3+**: Chen et al., "Encoder-Decoder with Atrous Separable Convolution," ECCV 2018
- **SAM**: Kirillov et al., "Segment Anything," ICCV 2023
- **MobileNetV3**: Howard et al., "Searching for MobileNetV3," ICCV 2019

Code and Datasets

Official Implementations:

- SAM: <https://github.com/facebookresearch/segment-anything>
- torchvision FCN: <https://pytorch.org/vision/stable/models.html>
- Our Mini-SAM: See artifact code provided earlier

Datasets:

- PASCAL VOC 2012: <http://host.robots.ox.ac.uk/pascal/VOC/voc2012/>
- Create subset: `random.sample(train_data, 500)`

Help

- Office hours: by appointment
- Course forum for questions