

An Intuitive Walkthrough of R-CNN: Understanding Object Detection Through Code

Let me walk you through this R-CNN implementation in a way that reveals the elegance and challenges of this foundational object detection approach. R-CNN, which stands for Regions with CNN features, was a breakthrough in 2014 because it showed how to effectively combine traditional computer vision with deep learning for object detection.

The Core Philosophy: Breaking Detection into Manageable Pieces

The fundamental insight of R-CNN is that object detection is really hard to solve end-to-end, so let's break it into separate, more tractable problems. Think of it like this: instead of asking a neural network to magically tell us "where are all the objects and what are they?" we break it down into three distinct questions: "where might objects be?", "what features define these regions?", and "how can we classify and refine these detections?"

The architecture starts with region proposals. In the code, you'll see two approaches implemented: `SelectiveSearchProposals` which uses OpenCV's selective search algorithm, and `SimplifiedProposals` which uses a multi-scale sliding window approach. Selective search is the classic R-CNN approach - it's a segmentation-based algorithm that groups pixels together based on color, texture, and size similarity to generate about 2,000 region proposals per image. The simplified version is there for practical purposes when selective search isn't available, generating proposals by sliding windows of various scales and aspect ratios across the image. The key point here is that these proposals are class-agnostic - we're just saying "hey, there might be something interesting here" without knowing what it is.

Feature Extraction: The CNN Backbone

Once we have our region proposals, we need to extract meaningful features from each region. This is where the `CNNFeatureExtractor` class comes in, using a pretrained ResNet-50 as the backbone. Here's what's happening: for each proposed region, we crop that rectangular patch from the image, warp it to a fixed size of 224×224 pixels (regardless of its original dimensions), and pass it through the CNN to get a 2048-dimensional feature vector. This warping is crucial - the CNN expects fixed-size inputs, so whether we're looking at a tiny 50×50 region or a large 200×300 region, we squeeze or stretch it to 224×224 .

The CNN itself has been modified by removing the final classification layer. We don't want it to classify ImageNet categories; we want those rich intermediate representations that

capture edges, textures, parts, and object-level features. The code uses the pretrained ResNet-50 weights, which means the network already knows how to extract useful visual features from images because it was trained on millions of ImageNet images. This transfer learning is fundamental to R-CNN's success - we're leveraging knowledge from one task (ImageNet classification) to help with another (object detection).

The Three-Stage Training Pipeline: Why R-CNN is Complex

Now here's where R-CNN gets interesting and frankly, a bit messy. The training happens in three completely separate stages, and this is both a strength and a weakness of the approach. Let's walk through each stage as implemented in the code.

Stage 1 is about fine-tuning the CNN for detection rather than classification. The code looks at each region proposal and asks: "does this overlap significantly with a ground truth object?" If a proposal has an Intersection over Union (IoU) of 0.5 or more with any ground truth box, it's labeled as positive. If it has IoU less than 0.1, it's negative. Everything in between is ignored as ambiguous. The goal is to adapt the pretrained CNN to be good at recognizing "objectness" - the quality of containing an object versus being background. In this implementation, we're actually skipping the fine-tuning step for simplicity and just using pretrained features, but in the original paper, you would retrain the CNN layers with these positive and negative examples.

Stage 2 trains the SVM classifiers, and this is where class-specific detection happens. Here's the subtle but important point: we train one binary SVM per class. Each SVM learns to distinguish "is this region a car?" or "is this region a person?" from background. The code extracts the CNN features for each proposal, then labels it positive if it overlaps with $\text{IoU} > 0.3$ with a ground truth box of that class. Notice the threshold changed from 0.5 to 0.3 - this is intentional! The SVM stage uses a more lenient threshold to get more training examples. For each class, we normalize the features using StandardScaler and train a linear SVM. The linear SVM is fast and works well in the high-dimensional feature space that the CNN provides.

Stage 3 is bounding box regression, which refines the initial proposals. Here's the idea: even if our region proposal roughly overlaps with an object, it's probably not perfectly aligned. The bounding box regressor learns a transformation to adjust the proposal's position and size to better match the ground truth. The code implements this using Ridge regression (regularized linear regression) to predict four transformation parameters: shifts in x and y coordinates, and scale changes in width and height. These transformations are computed in a clever way - instead of predicting absolute coordinates, we predict relative changes normalized by the proposal box dimensions. This makes the regression task easier and more generalizable.

Detection: Putting It All Together

When we run detection on a new image with the `detect` method, watch how all these pieces come together. First, we generate region proposals - maybe 2,000 of them. For each

proposal, we extract CNN features, giving us a 2048-dimensional vector. Then we run this feature vector through each of our class-specific SVMs. If any SVM fires with confidence above our threshold, we consider that a detection. But we don't just return the original proposal box - we pass the features and the proposal box through the trained regressor for that class, which refines the box to better fit the object.

The final step is Non-Maximum Suppression (NMS), which solves the problem of multiple overlapping detections. The code implements this per-class: for each class, we sort detections by confidence score, keep the highest scoring one, and remove any other detections that significantly overlap with it ($\text{IoU} > 0.3$). We repeat this process until no overlapping detections remain. This ensures we don't report the same object multiple times.

The Beautiful Inefficiency

Here's what makes R-CNN both elegant and impractical: it's incredibly modular and conceptually clean, but also painfully slow. For every single image, we're running the CNN forward pass up to 2,000 times - once for each proposal! Each forward pass warps the region, preprocesses it, and computes features. Then we have to run multiple SVM classifications and bounding box regressions. This is why R-CNN takes about 47 seconds per image on a GPU. The training is even more complex because it's this three-stage pipeline where you can't train everything end-to-end.

But understanding R-CNN deeply is crucial because all the modern detectors - Fast R-CNN, Faster R-CNN, and beyond - are essentially optimizations and improvements on these core ideas. They ask: how can we share computation? How can we make this differentiable end-to-end? How can we get rid of selective search? But the fundamental recipe - propose regions, extract features, classify, and refine - that's still there, just executed much more efficiently. That's why we start with R-CNN: it teaches us what object detection really means before we learn how to do it fast.

A Detailed Class-by-Class Walkthrough of R-CNN

Let me take you through each component of this R-CNN implementation, explaining how each class and its methods work together to create a complete object detection system.

The SelectiveSearchProposals Class: Finding Candidate Regions

The first class we encounter is `SelectiveSearchProposals`, which wraps OpenCV's selective search algorithm. The philosophy here is that we want to avoid the computational nightmare of checking every possible rectangular region in an image. Think about it - on a 400×400 image, there are millions of possible rectangles, and most of them are nonsense (half a chair, a piece of background, random noise).

The `generate_proposals` method is where the magic happens. It takes an image and asks selective search to find regions that might contain objects. Selective search works by starting with over-segmentation - breaking the image into many small regions based on color and texture - then hierarchically merging similar regions. The algorithm uses multiple similarity metrics: color similarity, texture similarity, size (prefer merging small regions first), and shape compatibility (prefer merging regions that fit well together). This produces a hierarchy of regions, and selective search outputs around 2,000 of these at various levels of the hierarchy.

The parameters matter here: `scale` controls the initial segmentation granularity, `sigma` affects the Gaussian smoothing applied before segmentation, and `min_size` filters out tiny proposals that are likely noise. The method processes these rectangles, converts them to the [x, y, width, height] format, and filters based on minimum size. The beauty of selective search is that it's class-agnostic and relatively fast - it uses traditional computer vision, not deep learning, so we're not burning GPU cycles here.

The SimplifiedProposals Class: A Practical Alternative

The `SimplifiedProposals` class exists because selective search requires specific OpenCV modules that aren't always available. This class implements a multi-scale sliding window approach - a classic computer vision technique, but done intelligently. Instead of a dense grid at one scale, we use multiple scales and aspect ratios.

In the `generate_proposals` method, we iterate over different scales (0.3, 0.5, 0.7 of the image dimensions) and aspect ratios (0.5, 1.0, 2.0 - capturing tall, square, and wide objects). For each scale-ratio combination, we create a window of the appropriate size and slide it across the image with a stride that's proportional to the window size. This is clever because smaller windows get smaller strides, maintaining roughly constant spatial coverage. The method generates potentially thousands of proposals, shuffles them randomly (to avoid spatial bias), and returns the top 2,000. It's not as sophisticated as selective search - we're just hoping that somewhere in this grid of multi-scale rectangles, we'll capture most objects reasonably well.

The CNNFeatureExtractor Class: Deep Features for Recognition

Now we get to the heart of R-CNN's representational power: the `CNNFeatureExtractor` class. This class wraps a pretrained ResNet-50 and adapts it for feature extraction rather than classification. The initialization is straightforward but important - we load ResNet-50

with ImageNet weights, then surgically remove the final fully connected layer. Why? Because that final layer is specialized for 1,000 ImageNet categories, but we want the rich 2048-dimensional feature representation that comes right before it.

The `transform` pipeline in the initialization defines how we'll preprocess image regions. Every region, regardless of its original size, gets resized to 224×224 pixels. This is a fundamental requirement of CNNs - they need fixed-size inputs because the fully connected layers (or in our case, the global average pooling) expect specific dimensions. We then normalize using ImageNet statistics (those specific mean and std values). This normalization is crucial - it ensures our image regions have the same distribution as the ImageNet images the network was trained on.

The `extract_region_features` method is called thousands of times during both training and inference, so understanding it is critical. It takes a full image and a bounding box [x, y, w, h], then crops out that region. Notice the defensive programming here - we clamp coordinates to be non-negative and ensure width and height are at least 1. If something goes wrong and we get an empty region, we return zeros rather than crashing. For valid regions, we apply the transform (which warps the region to 224×224 and normalizes it), add a batch dimension with `unsqueeze(0)` because PyTorch expects batches, then run it through the feature extractor. The `torch.no_grad()` context is important - we're not training the CNN here, just extracting features, so we don't need to compute gradients. The output is squeezed to remove the batch dimension, giving us our 2048-dimensional feature vector that captures everything the CNN knows about this region.

The BoundingBoxRegressor Class: Refining Detections

The `BoundingBoxRegressor` class solves an important problem: our initial region proposals are crude approximations. Even if a proposal overlaps with an object, it might be shifted or scaled incorrectly. This class learns to predict corrections that transform a proposal box into a tighter fit around the actual object.

The `compute_regression_targets` method implements the mathematical formulation from the R-CNN paper. Given a proposal box and a ground truth box, we compute four transformation parameters. For the x and y translations (t_x and t_y), we calculate the shift needed and normalize by the proposal width and height respectively. This normalization is key - it makes the targets scale-invariant. A 10-pixel shift means something very different for a 30-pixel wide box versus a 300-pixel wide box. For width and height transformations (t_w and t_h), we use logarithms of the scale factors. Why logarithms? Because they convert multiplication into addition, making the regression task more linear. If ground truth width is twice the proposal width, $\log(2) \approx 0.69$. If it's half, $\log(0.5) \approx -0.69$. Symmetric and nice for regression.

The `train_regressor` method trains a separate Ridge regression model for each class. This separation is important because different object categories might need different types of corrections - cars might typically need more width adjustment, people might need more height adjustment. For each training example, we compute the regression targets, normalize the features using StandardScaler (again, making the regression better conditioned), and fit

a Ridge regressor with strong regularization ($\text{alpha}=1000$). Ridge regression is chosen over ordinary least squares because it prevents overfitting in this high-dimensional feature space.

The `predict` method applies the learned transformation. Given features and a proposal box for a specific class, we first check if we even have a regressor for that class (if not, we just return the original proposal). We normalize the features the same way we did during training, predict the four transformation parameters, then apply them in reverse to get the refined box. We add the scaled translations to get new x and y, and we exponentiate the log-scale factors and multiply to get new width and height. This gives us a box that should fit the object much better than the original proposal.

The RCNN Class: Orchestrating the Detection Pipeline

The `RCNN` class is where everything comes together. It manages all the components and implements the complete training and inference pipeline. The initialization sets up the proposal generator, the CNN feature extractor, and creates dictionaries to hold SVMs and bounding box regressors for each class. The use of dictionaries indexed by `class_id` is elegant - it allows us to train models for only the classes we care about and easily add new classes later.

The `compute_iou` method is a fundamental building block used throughout the code. Intersection over Union is the standard metric for measuring bounding box overlap. The method converts boxes from $[x, y, w, h]$ to $[x_1, y_1, x_2, y_2]$ format, computes the intersection rectangle by finding the maximum of the top-left coordinates and minimum of the bottom-right coordinates, calculates the intersection area (handling the case where boxes don't overlap by using $\max(0, \dots)$), then divides by the union area. IoU ranges from 0 (no overlap) to 1 (perfect overlap), and we use different thresholds for different purposes throughout the pipeline.

Stage 1: Fine-tuning the CNN

The `train_stage1_finetune_cnn` method implements the first stage of training. For each training image, we generate proposals and compare them against ground truth boxes. The labeling strategy is carefully designed: proposals with $\text{IoU} \geq 0.5$ are positive examples (they contain an object well enough), proposals with $\text{IoU} < 0.1$ are negative examples (they're mostly background), and everything in between is discarded as ambiguous. This creates a clean binary classification problem.

The class balancing is crucial - we limit negative regions to match the number of positive regions. Without this, we'd be overwhelmed by negatives (most proposals don't overlap with objects), and the CNN would learn to just predict "background" for everything. The implementation here is simplified - in the original R-CNN paper, you would actually fine-tune the CNN weights using these positive and negative regions with a binary classification loss. But since fine-tuning a ResNet-50 takes significant time and data, this code uses the pretrained features directly, which still works reasonably well thanks to transfer learning.

Stage 2: Training the SVMs

The `train_stage2_svm` method is where we learn class-specific detection. Notice we loop over each class independently - this is one SVM per class, not one multi-class SVM. For each class, we generate proposals on all training images and extract features for each proposal. The labeling here uses a lower IoU threshold (0.3) compared to Stage 1 (0.5). Why? The paper found that being more permissive at this stage gives better performance - it provides more positive training examples and allows the SVM to learn a slightly broader notion of what constitutes that object class.

For each proposal, we check if it overlaps with $\text{IoU} > 0.3$ with any ground truth box of the current class. If yes, it's labeled 1 (positive for this class); otherwise, it's labeled 0 (negative/background). We collect all features and labels, normalize with StandardScaler, and train a linear SVM. The LinearSVC with $C=0.01$ provides regularization to prevent overfitting. The `dual='auto'` parameter lets sklearn choose the optimization algorithm based on the number of samples. Linear SVMs work well here because the CNN features are already highly discriminative - we're operating in a space where classes are roughly linearly separable.

Stage 3: Training Bounding Box Regressors

The `train_stage3_bbox_regression` method fine-tunes our detections. Again, we loop over classes independently. For each class, we look for proposals that tightly overlap ($\text{IoU} > 0.6$) with ground truth boxes of that class. This higher threshold (compared to 0.3 in Stage 2) ensures we're learning to refine already-good proposals, not trying to fix terrible ones. For each good pair of (proposal, ground truth), we extract features and add them to the training set. If we find sufficient training examples, we train the regressor for this class. Some classes might not have enough training data, especially if they appear rarely - the code handles this gracefully by skipping regressor training for those classes.

The Complete Training Pipeline

The `train` method simply orchestrates the three stages in sequence. This is the high-level interface - you pass in images and annotations, and it handles all the complexity internally. The three-stage nature is a key characteristic of R-CNN: we can't train everything jointly end-to-end like modern detectors. Each stage depends on the previous one, but they optimize different objectives with different algorithms (SGD for CNN, SVM optimization for classifiers, Ridge regression for bounding boxes).

Detection: Running Inference

The `detect` method is where we use our trained model to find objects in new images. First, we generate proposals - the same process as training, giving us up to 2,000 candidate regions. For each proposal, we extract CNN features. Now comes the classification step: we run each proposal's features through all class SVMs. The SVM's `decision_function` gives us a confidence score (not a probability - it's the distance from the separating hyperplane). If this score exceeds our threshold for any class, we have a detection. But we don't return the raw proposal box - we pass it through the bounding box regressor for that class to get a refined box.

At this point, we might have hundreds of detections, many overlapping. That's where NMS comes in.

Non-Maximum Suppression

The `apply_nms` method implements a classic algorithm to eliminate redundant detections. We group detections by class (NMS is applied per-class because a car and a person might legitimately overlap). For each class, we sort detections by confidence score descending. We then greedily select the highest-scoring detection, add it to our keep list, and remove all other detections that significantly overlap with it ($\text{IoU} > \text{threshold}$). We repeat until no detections remain. This ensures that each object is detected at most once - we keep the most confident detection and suppress the others.

The implementation uses a simple while loop with nested filtering. For the best detection, we create a new list containing only detections that don't overlap too much with it, then continue with this filtered list. This is $O(n^2)$ in the number of detections, but it's simple and correct.

Saving and Loading Models

The `save_model` and `load_model` methods handle persistence. Notice what we save: the SVMs, their scalers, the bounding box regressors, and the number of classes. We don't save the CNN because it's just the pretrained ResNet-50 - we can reload that anytime. We use pickle for serialization, which works well for sklearn models and our custom classes. This allows you to train once and deploy the trained detector elsewhere.

The Visualization Function: Seeing the Results

The `visualize_detections` function isn't part of the core algorithm, but it's essential for understanding what's happening. It creates a matplotlib figure, draws the image, then overlays each detection as a colored rectangle with a label. The label includes the class name (or ID) and confidence score. Using different colors for different classes helps visually distinguish what's been detected. The bounding boxes are drawn with colored borders, and the text labels have colored backgrounds matching their boxes for easy association.

The Demo: Putting It All Together

The main block demonstrates the complete pipeline with synthetic data. The `create_synthetic_image` function generates images with colored rectangles at random positions - a simple but effective way to test the detector without requiring real datasets. Each rectangle is a different color and gets assigned to one of three classes.

The demo creates five training images, initializes an R-CNN with three classes, trains it through all three stages, then tests on a new image. By setting the confidence threshold very low (0.0) during testing, we can see what the detector finds even with low confidence, which is useful for debugging. The detections are printed and visualized, giving us a complete view of what the system learned.

This walkthrough reveals R-CNN's structure: it's modular, with each component solving a specific subproblem. The proposal generator finds candidates, the CNN extracts features, the SVMs classify, the regressors refine, and NMS cleans up. Understanding each piece and how they interact gives you the foundation to understand all the improvements that came after - Fast R-CNN, Faster R-CNN, and modern detectors.

A Detailed Class-by-Class Walkthrough of Fast R-CNN

Let me guide you through the Fast R-CNN implementation, revealing how it solves R-CNN's fundamental inefficiency through elegant architectural innovations. Fast R-CNN was a breakthrough in 2015 because it recognized that the biggest bottleneck in R-CNN wasn't the algorithm itself - it was the wasteful computation of running the CNN 2,000 times per image.

The Core Innovation: From Sequential to Shared Computation

Before diving into the classes, understand the paradigm shift: R-CNN processes each region proposal independently through the CNN, which means for an image with 2,000 proposals, we do 2,000 complete forward passes through ResNet-50. Fast R-CNN flips this around - we run the CNN once on the entire image to get a feature map, then extract features for each proposal from this shared representation. This single change provides a roughly 2,000× reduction in CNN computation.

The RoIPooling Class: Bridging Variable and Fixed Sizes

The `RoIPooling` class is the technical heart of Fast R-CNN, solving a fundamental problem: how do we extract fixed-size feature representations from variable-size regions on a feature map? This is crucial because the detection head expects consistent input dimensions, but our region proposals come in all shapes and sizes.

The initialization is simple - we just specify the output size, typically 7×7 , which will be the spatial dimensions of our pooled features. This choice isn't arbitrary; it's a trade-off between spatial resolution (higher is more detailed) and computational cost (larger means more parameters in subsequent layers).

The **forward** method is where the sophisticated geometry happens. It receives three inputs: the shared feature map from the CNN backbone (say, $2048 \times 14 \times 14$ for a 224×224 input image), a batch of Rols specified in image coordinates, and the original image size. The first critical step is computing the spatial scale - the ratio between feature map dimensions and image dimensions. If our image is 224×224 and our feature map is 14×14 , the spatial scale is $14/224 \approx 0.0625$. This tells us how to map coordinates from image space to feature space.

For each RoI, we extract its batch index (which image in the batch it belongs to) and its coordinates $[x_1, y_1, x_2, y_2]$ in image space. We multiply these coordinates by the spatial scale to get feature map coordinates. Here's where the pooling operation gets interesting: we divide the RoI on the feature map into a grid of output_size bins (7×7 in our case). For each bin, we compute its boundaries on the feature map, which might not align perfectly with pixel boundaries - a bin might span from pixel 2.3 to pixel 4.8, for example.

The pooling then applies max pooling within each bin. We use adaptive max pooling to handle the irregular bin sizes gracefully. If a bin maps to feature map coordinates [2.3, 4.8], we take the integer bounds [2, 4] and max pool that region to a single value. This process repeats for all channels (2048 of them) and all spatial locations (7×7), giving us a fixed $2048 \times 7 \times 7$ output regardless of the input RoI size. A 50×50 proposal and a 200×200 proposal both become $2048 \times 7 \times 7$ feature tensors, ready for the detection head.

The defensive programming throughout - clamping coordinates, handling edge cases, ensuring bins have positive area - isn't just good practice. It prevents subtle bugs that would cause training to crash when proposals extend beyond image boundaries or have degenerate geometries.

The problem RoI Pooling solves

After your backbone CNN (e.g., ResNet), you have a **feature map** that is smaller than the original image — say **14×14** instead of 224×224 .

Now, the **region proposals** (Rois) are boxes in the original image coordinates, like:

```
Object 1 → [x1, y1, x2, y2] = [50, 30, 180, 200]  
Object 2 → [x1, y1, x2, y2] = [100, 60, 130, 100]
```

These boxes are different sizes — one might cover a big object, the other a small one.

But the **Fast R-CNN head** (the fully connected layers that classify + regress) expects a *fixed-size vector input* — like **2048×7×7** — for every RoI.

So RoI Pooling's job is:

Turn each arbitrary-sized region into a fixed-size 7×7 feature map.

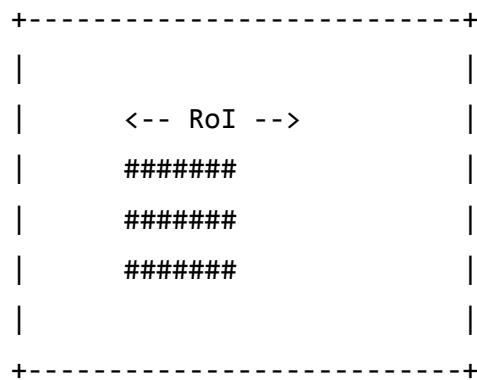
2The 7×7 window does *not* move over the image

This is key:

- The 7×7 grid is *not a sliding filter* like in convolution.
- It's a **conceptual grid** that divides one *proposal box* into 7×7 **bins**.

Think of it like this:

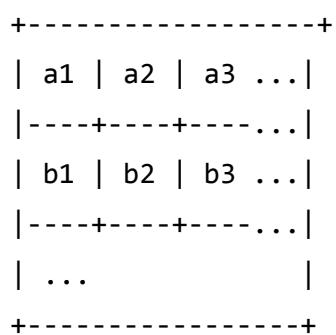
Feature Map (14×14)



Now, the region with # corresponds to the projection of one RoI from the original image (scaled down to feature map space).

RoI Pooling divides that region only into a 7×7 grid:

RoI Region (on feature map)



Each small bin (like a1) will be pooled (max/avg) into a single number per channel. So after processing all 7×7 bins, we end up with a **7×7×C** tensor.

That 7×7 tensor *represents the whole RoI*, not the entire image.

◆ 3What happens step-by-step (from your code)

Let's map that to the code:

```
for roi_idx in range(num_rois):
    # 1. Extract the region of interest
    x1_feat = int(x1 * spatial_scale_w)
    y1_feat = int(y1 * spatial_scale_h)
    x2_feat = int(x2 * spatial_scale_w)
    y2_feat = int(y2 * spatial_scale_h)
```

Now we have **the coordinates of the RoI on the feature map** (not the whole image).

Then:

```
# 2. Divide that region into a 7x7 grid of bins
bin_h = roi_height / output_h # roi_height / 7
bin_w = roi_width / output_w # roi_width / 7
```

Then for each bin:

```
# 3. Take the corresponding feature map slice
roi_bin = feature_map[batch_idx, :, y_start:y_end, x_start:x_end]

# 4. Pool it to a single value (max)
pooled_features[roi_idx, :, i, j] = F.adaptive_max_pool2d(roi_bin, (1,1))
```

That produces a **single value per bin per channel**.

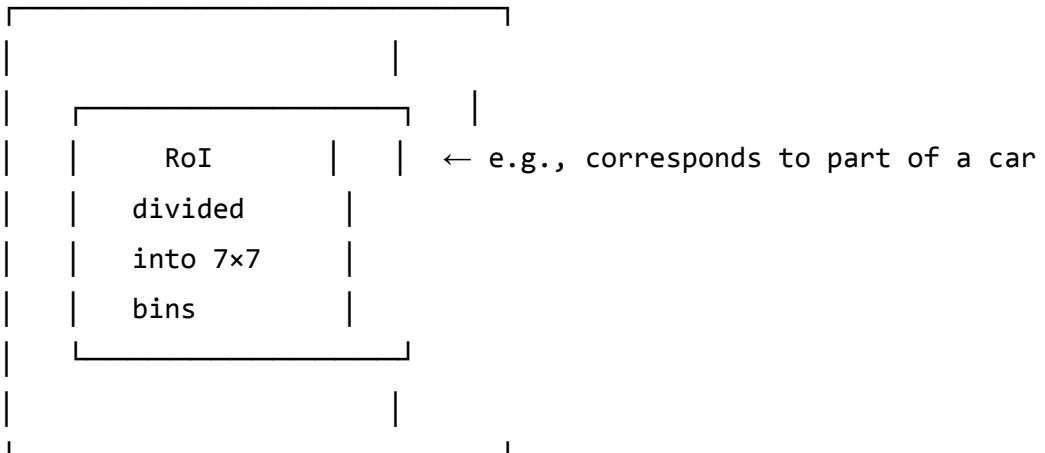
After looping over all 7×7 bins, you get:

```
[  
 \text{Output shape: } [C, 7, 7]  
 ]
```

4Intuitive picture

Let's visualize the idea conceptually:

Feature map (14×14)



Each bin → pooled to 1 value → forms 7×7 representation of that RoI.

That's why **the 7×7 “window” doesn't slide**;
it's redefined *for each proposal box individually*.

◆ 5 Why 7×7?

It's a design compromise:

- Larger → more detail, but more parameters for FC layers.
- Smaller → less spatial information, less precision.

7×7 was empirically found to give a good balance in the original Fast R-CNN.

Summary

Concept	Meaning
Input	Shared feature map (e.g., 2048×14×14)

Rois	Variable-size boxes in image space
Roi Pooling	Projects each RoI $\rightarrow 7 \times 7$ grid on feature map
Operation	Max-pools each bin \rightarrow single value
Output	Fixed-size $2048 \times 7 \times 7$ tensor per RoI
Purpose	Makes detection head possible (fixed-size input)

Intuitive analogy:

Roi Pooling “zooms in” on each region proposal, samples its corresponding patch on the feature map, and “compresses” it into a fixed 7×7 representation — like resizing each detected part of the image to a common scale before passing it to the classifier.

Would you like me to make this into a **two-slide Beamer pair** (concept + diagram/code summary) for your *Fast R-CNN section*?

It would include a small $14 \times 14 \rightarrow 7 \times 7$ TikZ grid illustration to help students visualize how Roi pooling works geometrically.

The FastRCNNHead Class: Multi-Task Prediction

The `FastRCNNHead` class implements the detection head that sits atop the Roi pooled features. This is where Fast R-CNN makes another key improvement over R-CNN: instead of training separate SVMs for classification and separate regressors for bounding boxes, we train both tasks jointly in a single network with a multi-task loss.

The initialization constructs a simple but effective architecture. We have two fully connected layers (`fc6` and `fc7`), each with 4096 units. These are followed by two separate output branches: `cls_score` for classification and `bbox_pred` for bounding box regression. The classification head outputs scores for `num_classes + 1` categories (the `+1` is for background). The bounding box regression head outputs 4 coordinates for each class - that's why it has $4 \times (\text{num_classes} + 1)$ outputs. We predict class-specific box refinements because different object types might need different transformations.

The weight initialization uses Gaussian initialization with small standard deviation (0.01) for all layers. This is important for training stability - starting with weights too large causes gradients to explode, too small and learning is painfully slow.

The `forward` method implements the computation graph. The RoI pooled features, already flattened to a vector, pass through fc6 with ReLU activation and dropout (0.5 probability during training). Dropout is crucial here - with 4096 units, these layers have massive capacity and would easily overfit without regularization. The same pattern repeats for fc7. Then we branch: one path produces classification scores (logits, not probabilities yet), another produces bounding box deltas. These are returned as separate tensors, allowing the loss function to handle them appropriately.

The FastRCNN Class: The Complete Architecture

The `FastRCNN` class orchestrates the entire detection pipeline, bringing together the backbone, RoI pooling, and detection head into a unified model. This is the main class you instantiate and train.

The initialization sets up the three-part architecture. First, we load a pretrained ResNet-50 and extract everything except the final average pooling and classification layer. We want the convolutional feature maps, not the global image classification. This gives us a backbone that outputs 2048-channel feature maps. The spatial dimensions depend on the input image size - for a 224×224 input, we get roughly 14×14 features after all the pooling and striding in ResNet.

We then instantiate our RoI pooling layer with the specified output size (default 7×7). The detection head gets initialized with the correct input dimension: backbone channels (2048) times the pooling output size (7×7) equals $2048 \times 7 \times 7 = 100,352$ features per RoI. This is why we need those fc layers to compress this high-dimensional representation into something manageable.

The optional freezing of early backbone layers is a practical choice. The first 30 or so parameters (early convolutional filters) learn very general features - edges, colors, textures - that don't need much adaptation for detection. Freezing them speeds up training and reduces overfitting.

The `forward` method reveals the elegant simplicity of Fast R-CNN's design. We receive a batch of images and proposals. First - and this is the key innovation - we run the entire image batch through the backbone CNN once. This single forward pass produces feature maps for all images. Then, for each proposal, we use RoI pooling to extract its features from the appropriate feature map (the `batch_idx` in the proposal tells us which image). We flatten these pooled features and pass them through the detection head, which outputs classification scores and bounding box deltas for all proposals simultaneously.

Compare this to R-CNN: instead of cropping each proposal, resizing it, and running it through the CNN (2,000 times), we run the CNN once and use RoI pooling to extract proposal features. The computational savings are enormous.

The Smooth L1 Loss Function: Robust Regression

The `smooth_l1_loss` function implements a crucial component of the multi-task loss. This is also called Huber loss, and it combines the best properties of L1 and L2 losses. When the prediction error is small (less than beta, typically 1.0), we use L2 loss: $0.5 \times x^2$. This provides smooth gradients near the optimum, helping convergence. When the error is large, we use L1 loss: $|x| - 0.5$. This is less sensitive to outliers than L2, which squares the error and thus heavily penalizes large mistakes.

Why does this matter? In bounding box regression, we occasionally get bad proposals that are far from any object. With L2 loss, these outliers would dominate the gradient and destabilize training. With L1 loss alone, gradients are constant regardless of error magnitude, which slows convergence. Smooth L1 gives us stable gradients for small errors and robustness to outliers for large errors.

The implementation uses `torch.where` to select between the two formulas based on the threshold. The sum reduction means we're adding up losses across all coordinates of all boxes, which the calling function will normalize appropriately.

The FastRCNNLoss Class: Multi-Task Learning

The `FastRCNNLoss` class implements the joint optimization objective that defines Fast R-CNN's training. The philosophy is that classification and localization are related tasks that should be learned together, not separately as in R-CNN.

The initialization just stores the number of classes and a weighting factor `lambda_bbox` (default 1.0) that balances the two loss terms. This balance is important - if classification loss dominates, the network ignores localization; if regression dominates, classification suffers.

The `__call__` method computes the complete loss. For classification, we use standard cross-entropy loss comparing predicted scores against ground truth labels. Label 0 is background, labels 1 through `num_classes` are object categories. This is straightforward and identical to any classification problem.

The bounding box regression loss is more subtle. We receive `bbox_deltas` shaped $[N, 4 \times (\text{num_classes} + 1)]$ - that's four coordinates for each possible class. But we only want to apply regression loss for the true class of each proposal. Background proposals (label 0) shouldn't contribute to regression loss at all. We first reshape the deltas to $[N, \text{num_classes} + 1, 4]$ to separate the per-class predictions. Then we use advanced indexing with the ground truth labels to select the relevant deltas:

`bbox_deltas[torch.arange(batch_size), labels]` picks out the prediction for the true class of each proposal.

The `bbox_inside_weights` tensor is binary: 1 for foreground objects (which should regress to their ground truth), 0 for background (which we ignore). We multiply both predictions and targets by these weights before computing smooth L1 loss. This effectively zeros out the loss for background proposals. We normalize by batch size to make the loss magnitude comparable to classification loss.

The final loss is simply the sum: $L = L_{cls} + \lambda \times L_{bbox}$. Both tasks contribute to the gradient, and the network learns to balance them. This joint training is faster and often more accurate than R-CNN's separate training stages.

The SimplifiedProposalGenerator Class: Proposal Creation

The `SimplifiedProposalGenerator` class provides region proposals for training and testing. In a real Fast R-CNN system, you'd use Selective Search or later, a Region Proposal Network. This simplified version generates random boxes, which is sufficient for demonstration and testing.

The `generate_proposals` method is straightforward: it creates random rectangles on the image. For each proposal, we pick a random top-left corner, ensuring we leave room for the box (hence $w - 20$ as the maximum $x1$). Then we pick a bottom-right corner, ensuring the box has minimum size ($x1 + 20$) and doesn't exceed the image ($x1 + 100$ capped at width). This gives us boxes of varying sizes and aspect ratios, roughly simulating what a real proposal generator would produce.

The randomness is actually useful during training - it ensures we see diverse proposals and the network doesn't overfit to specific proposal patterns. In practice, you'd want better proposals (higher recall of actual objects), but for understanding Fast R-CNN's mechanics, this suffices.

The FastRCNNTrainer Class: Training Pipeline

The `FastRCNNTrainer` class manages the complete training process, handling data preparation, forward passes, loss computation, and optimization. This is where we see how all the pieces fit together.

The initialization sets up the model on the appropriate device (GPU if available), initializes the loss function, and creates an SGD optimizer with momentum and weight decay. These hyperparameters ($lr=0.001$, $momentum=0.9$, $weight_decay=0.0005$) are standard for training detection networks and provide a good balance between convergence speed and stability.

The `compute_iou` method is identical in concept to R-CNN's version but works with $[x1, y1, x2, y2]$ format directly. It's used throughout training to assign labels to proposals based on overlap with ground truth.

The `compute_bbox_targets` method calculates regression targets, but notice the formulation is slightly different from R-CNN. Here we work with box centers and dimensions explicitly. We convert each box to center coordinates (px, py) and dimensions (pw, ph), then compute the transformation needed to move the proposal center to the ground truth center (normalized by dimensions) and the log-scale factors for size changes. This parameterization makes the regression task more symmetric and easier to learn.

The `prepare_training_data` method is crucial - it converts the raw annotations into the tensors needed for training. For each image and its proposals, we find the best matching ground truth box for each proposal using IoU. If $\text{IoU} \geq 0.5$, the proposal is labeled with the object's class (shifted by +1 because 0 is background) and we compute regression targets to that ground truth box. The `bbox_weight` is set to ones, indicating this proposal should contribute to regression loss.

If $\text{IoU} < 0.5$, the proposal is background (label 0), with zero regression targets and zero weights. This proposal won't contribute to regression loss, only to classification loss. The method accumulates all proposals across the batch into tensors, with each RoI stored as `[batch_idx, x1, y1, x2, y2]` so RoI pooling knows which image it belongs to.

The `train_epoch` method runs one complete training iteration over the provided images. It generates proposals for each image, converts images to tensors with proper normalization (dividing by 255 to get $[0, 1]$ range), and prepares all the training data. The forward pass is beautifully simple: images and RoIs go into the model, classification scores and bbox deltas come out. The loss function computes the multi-task loss, and we do standard backpropagation: zero gradients, backward pass, optimizer step.

The return values (`total_loss`, `cls_loss`, `bbox_loss`) let us monitor training progress and understand how each task contributes to the overall objective.

The Helper Functions: Demonstration and Visualization

The `create_synthetic_dataset` function generates simple training data for demonstration. Each image contains 2-3 randomly placed colored rectangles representing objects of different classes. This simplicity is actually pedagogically valuable - it lets students see that the algorithm works without the complexity of real images obscuring the mechanics.

The `demo_fast_rcnn` function provides a complete, runnable demonstration that showcases Fast R-CNN's advantages. It creates synthetic data, initializes the model, trains for 10 epochs, and critically, shows the computational comparison with R-CNN. The key insight is printed clearly: R-CNN needs 10,000 CNN passes for 5 images with 2,000 proposals each, while Fast R-CNN needs just 5 CNN passes - one per image. This $\sim 2,000\times$ reduction in computation is the headline achievement.

The inference demonstration at the end shows how to use the trained model: generate proposals, run the model to get scores and deltas, apply softmax to get probabilities, and filter detections by confidence. This mirrors how you'd deploy Fast R-CNN in practice.

The Elegant Solution to R-CNN's Bottleneck

Understanding Fast R-CNN class-by-class reveals its elegant design philosophy. Every component serves a specific purpose: RoI pooling enables shared computation, the detection head unifies classification and regression, smooth L1 loss stabilizes training, and multi-task learning improves both tasks simultaneously. The architecture is end-to-end

trainable, unlike R-CNN's three-stage pipeline, which means we can backpropagate through everything and optimize the entire system jointly.

The computational savings come from a simple but powerful insight: convolutional features are translation-equivariant. When we compute features for the full image, features for each region are already implicitly there - we just need to extract them. RoI pooling is the tool that makes this extraction possible while maintaining the fixed-size requirement of fully connected layers.

This is why Fast R-CNN was such a breakthrough: it took R-CNN's fundamentally sound approach to detection (proposals + CNN features + classification + regression) and made it practical by eliminating the computational bottleneck. The next step, Faster R-CNN, would eliminate the remaining bottleneck - Selective Search - by making region proposals learnable too. But Fast R-CNN proved the core concept: shared computation is the key to efficient object detection.

A Detailed Class-by-Class Walkthrough of Anchor Generation and RPN

Let me guide you through the anchor generation mechanism and Region Proposal Network (RPN), which together represent the crucial innovation that transformed Faster R-CNN into a truly end-to-end learnable system. Understanding anchors deeply is fundamental because they appear in virtually all modern object detectors - from RetinaNet to YOLO to modern transformer-based detectors.

The Paradigm Shift: From Hand-Crafted to Learned Proposals

Before diving into the code, let's understand the conceptual leap. Fast R-CNN still relied on Selective Search, a hand-crafted algorithm that took about 2 seconds per image and couldn't be trained. The insight behind Faster R-CNN was profound: what if we could learn to

generate proposals using the same CNN features we're already computing? Instead of running a separate algorithm, we add a small network that looks at the CNN features and says "here's where objects might be." But there's a chicken-and-egg problem: how do you predict arbitrary bounding boxes from scratch? The answer is anchors - pre-defined reference boxes that the network learns to classify (object or background?) and refine (how should I adjust this box?).

The AnchorGenerator Class: Templates for Detection

The `AnchorGenerator` class implements the core concept of anchor boxes. Think of anchors as a vocabulary of box shapes - we're giving the network a set of templates and asking it to choose which ones look like objects and how to adjust them.

The initialization accepts three critical parameters that define our anchor vocabulary. The `scales` parameter (default [8, 16, 32]) controls the absolute sizes of our anchor templates in pixels. These numbers are multiplied by the stride, so with a stride of 16, we get actual anchors of size 128, 256, and 512 pixels. This multi-scale approach allows us to detect small, medium, and large objects without building an image pyramid. The `aspect_ratios` parameter (default [0.5, 1.0, 2.0]) defines the shape diversity - 0.5 gives tall/thin boxes (height is twice the width), 1.0 gives squares, and 2.0 gives wide/flat boxes (width is twice the height). The `stride` parameter (default 16) reflects how much the CNN downsampled the image - if our input is 224×224 and the feature map is 14×14, the stride is 16.

The beauty of this parameterization is its efficiency: with just 3 scales and 3 aspect ratios, we get 9 anchors per location, covering a wide range of object sizes and shapes. The `num_anchors` attribute simply multiplies these to give us the total number of templates per spatial location.

Generating Base Anchor Templates

The `generate_base_anchors` method creates the fundamental templates, all centered at the origin (0, 0). This is elegant because we can generate these templates once and then simply shift them to all spatial locations. The method iterates over each scale and aspect ratio combination. For each pair, we compute the anchor dimensions using a clever geometric approach: we want an anchor with $\text{area} = \text{scale}^2$ and aspect ratio = height/width. Solving these two equations, we get $\text{width} = \sqrt{(\text{area}/\text{ratio})}$ and $\text{height} = \text{width} \times \text{ratio}$. This ensures that anchors with the same scale have the same area, regardless of aspect ratio, which provides uniform coverage of the scale space.

The anchor coordinates are computed relative to the center point (0, 0). We calculate half-widths and half-heights, then construct the box as [-w/2, -h/2, w/2, h/2] in [x1, y1, x2, y2] format. The returned array has shape [num_anchors, 4], where num_anchors is 9 in the default configuration. These templates will be shifted to every location on the feature map.

Tiling Anchors Across the Image

The `generate_anchors` method is where we go from 9 base templates to potentially thousands of anchors covering the entire image. It takes the feature map size (e.g., 14×14) and the original image size (e.g., 224×224) and produces anchors at every spatial location. First, we generate our base templates centered at the origin. Then comes the tiling process: we iterate over every position (i, j) in the feature map. For each position, we compute its corresponding center in image coordinates. This mapping is crucial - if we're at feature map location (i, j) , the center in image space is $(j \times \text{stride} + \text{stride}/2, i \times \text{stride} + \text{stride}/2)$. The $\text{stride}/2$ offset ensures we're at the center of the receptive field, not the corner.

For each feature map location, we shift all base anchors to that center. If a base anchor is $[-4, -8, 4, 8]$ and our center is $(120, 120)$, the shifted anchor becomes $[116, 112, 124, 128]$. We repeat this for all 9 base anchors at this location, then move to the next location. The result is `feature_map_height \times feature_map_width \times num_anchors` total anchors. For a 14×14 feature map with 9 anchors per location, that's 1,764 anchors.

The final step is clipping anchors to image boundaries. Some anchors, especially near the edges, will extend beyond the image. We clamp all coordinates to $[0, \text{width}]$ and $[0, \text{height}]$ to ensure anchors stay within valid regions. This prevents the network from trying to predict objects that extend into non-existent image areas.

Visualizing Base Anchors

The `visualize_base_anchors` method creates an educational visualization showing all 9 base anchor templates centered at the origin. It generates a matplotlib figure with all templates drawn as colored rectangles. Each anchor gets a different color from a colormap, and we label it with its scale and aspect ratio (e.g., "S=16, R=1.0"). The center point is marked with a red star to emphasize that these are all centered at the same location.

The visualization uses equal aspect ratio and draws grid lines plus axis lines through the origin, making it clear that these are templates relative to a reference point. This figure is crucial for teaching because it shows students exactly what we mean by "multi-scale, multi-aspect-ratio templates" - they can see that some boxes are small and square, others are large and rectangular, covering a diverse range of object appearances.

Visualizing Anchors on Images

The `visualize_anchors_on_image` method creates a two-panel figure that shows how anchors tile across an actual image. The left panel focuses on one spatial location, drawing all 9 anchors at that point. This helps students understand that every location has the same 9 templates, just shifted in space. The right panel shows a sample of anchors across the entire image. To avoid visual clutter, we don't draw all 1,764 anchors - instead, we draw a grid of blue dots showing all anchor centers, then draw actual boxes at every other location for just one anchor type (the largest square anchor).

This visualization is pedagogically powerful because it bridges the gap between the abstract concept ("we have templates at every location") and the concrete reality ("here's what that actually looks like on an image"). Students can see the dense, regular tiling and understand

why this coverage ensures we won't miss objects - no matter where an object appears, there's likely an anchor nearby that roughly matches its size and shape.

The statistics printed after the visualization quantify what we're seeing: feature map size, anchors per location, total anchors, stride, scales, and aspect ratios. This helps students connect the parameters to the actual output.

The Demonstration Function: Teaching Anchor Concepts

The `demo_anchor_properties` function orchestrates a complete educational demonstration. It starts by creating an `AnchorGenerator` with the standard Faster R-CNN configuration (3 scales, 3 aspect ratios, stride 16). Then it walks through anchor generation step by step, printing the results in a formatted table. For each of the 9 base anchors, it shows the scale, aspect ratio, actual width and height, and area. This table reveals important patterns: anchors with the same scale have similar areas, and aspect ratios create the expected shape variations (tall boxes have height > width, wide boxes have width > height).

The function then demonstrates anchor properties with clear bullet points: multiple scales for size diversity, multiple shapes for aspect ratio diversity, ability to detect varied objects, and the fact that each location gets the same set of candidates. This repetition reinforces the core concepts. Next, it generates the full set of anchors for an image and prints statistics about the total coverage. Finally, it calls both visualization functions to create the figures.

The concluding explanation is critical for student understanding. It explicitly connects anchors to the key problems they solve: multi-scale detection without image pyramids, multi-shape detection for diverse objects, dense coverage so nothing is missed, and computational efficiency because anchors are pre-defined templates. The RPN just needs to answer two questions per anchor: "Is there an object here?" and "How should I adjust this box?" This is much simpler than predicting boxes from scratch.

The `AnchorGeneratorTorch` Class: PyTorch Integration

The `AnchorGeneratorTorch` class is a PyTorch module that implements the same anchor generation logic but in a differentiable, GPU-friendly way. It inherits from `nn.Module`, making it a proper component of the neural network. The initialization is similar to the NumPy version, storing scales, aspect ratios, and stride. But there's an important difference: it uses `register_buffer` to store the base anchors. This tells PyTorch that `base_anchors` is a tensor that should be moved to GPU along with the model, but it's not a learnable parameter - it's a fixed template.

The `_generate_base_anchors` method computes templates identically to the NumPy version, but returns a PyTorch tensor instead of a NumPy array. This ensures compatibility with the rest of the network - all intermediate computations will be on the same device (GPU if available) and in the same framework.

The `forward` method is where PyTorch's tensor operations shine. Instead of using nested loops like the NumPy version, it uses vectorized operations. It creates shift vectors for all spatial locations using `torch.arange` and `torch.meshgrid`. The meshgrid creates two 2D grids: one with x-coordinates, one with y-coordinates. We flatten these and stack them to get a $[H \times W, 4]$ tensor where each row is $[center_x, center_y, center_x, center_y]$ - the center point duplicated because we need to shift both x_1, y_1 , and x_2, y_2 .

The broadcasting operation is beautifully concise: `shifts.view(-1, 1, 4) + base_anchors.view(1, -1, 4)` adds every shift ($H \times W$ of them) to every base anchor (9 of them), producing $[H \times W, 9, 4]$ which we reshape to $[(H \times W \times 9), 4]$. This single line replaces all the nested loops from the NumPy version. The in-place clamping (`clamp_`) ensures memory efficiency. This PyTorch implementation is faster and integrates seamlessly with the rest of the network.

The RPN Class: Learning Which Anchors Are Objects

The `RPN` class is the heart of Faster R-CNN's proposal generation. It's a small neural network that looks at CNN features and predicts, for each anchor, whether it contains an object and how to adjust it. The architecture is intentionally simple to keep the proposal generation fast.

The initialization sets up three convolutional layers. The first is a 3×3 conv with 512 output channels - this provides a small amount of additional feature processing. The `padding=1` ensures the spatial dimensions don't change. Then we have two 1×1 convolutions that produce the actual predictions. The `cls_logits` layer outputs `num_anchors` channels (9 in the default case) - one objectness score per anchor. The `bbox_pred` layer outputs `num_anchors \times 4` channels (36 in the default case) - four regression values (tx, ty, tw, th) per anchor.

Why 1×1 convolutions for the prediction heads? Because we want independent predictions at each spatial location. A 1×1 conv is essentially a per-pixel fully connected layer - it looks at the features at location (i, j) and predicts values for that location only. This is perfect for anchor prediction because anchors at different locations are independent.

Weight initialization uses small random values (`std=0.01`) for all layers, with biases initialized to zero. Small initialization is important for stable training - large initial weights could cause extreme predictions that destabilize learning.

Forward Pass: Dense Predictions

The `forward` method implements the RPN's computation. It receives CNN features with shape $[B, C, H, W]$ and produces predictions for all anchors at all locations. First, the features pass through the 3×3 conv with ReLU activation. This intermediate representation is shared by both prediction heads - an efficient design because the same features inform both objectness and box refinement.

The objectness predictions come from the `cls_logits` layer, producing $[B, \text{num_anchors}, H, W]$. We permute this to $[B, H, W, \text{num_anchors}]$ because it's more intuitive to think of it as "for each location (H, W) , we have num_anchors scores." This also matches how we'll process them during training and inference.

The `bbox_deltas` predictions similarly go from $[B, \text{num_anchors} \times 4, H, W]$ to $[B, H, W, \text{num_anchors} \times 4]$. At location (i, j) , we have 36 values: 4 deltas for each of 9 anchors. These deltas will be used to transform the anchor boxes into refined proposals.

Notice that the RPN doesn't use any global information - predictions at each location depend only on local features. This spatial locality is a key property that enables efficient computation.

The `RPNLoss` Class: Training the Proposal Network

The `RPNLoss` class implements the loss function that trains the RPN. It needs to solve a challenging problem: given anchors and ground truth boxes, which anchors should we train as positive examples (containing objects) and which as negative examples (background)?

The initialization stores hyperparameters that control this assignment. The `pos_iou_threshold` (default 0.7) and `neg_iou_threshold` (default 0.3) define the IoU boundaries for labeling. The `batch_size` (default 256) limits how many anchors we use per image to prevent overwhelming the training with too many samples. The `pos_fraction` (default 0.5) aims to balance positive and negative examples, though in practice we often have fewer positives than desired.

Computing IoU Efficiently

The `compute_iou` method calculates intersection over union between two sets of boxes using vectorized operations. It receives `boxes1` $[N, 4]$ and `boxes2` $[M, 4]$ and produces an $[N, M]$ IoU matrix. The implementation computes areas for all boxes, then uses broadcasting to find intersections between all pairs. The key is `boxes1[:, None, :2]` which has shape $[N, 1, 2]$, and `boxes2[:, :, :2]` which has shape $[M, 2]$ - broadcasting these produces $[N, M, 2]$. This single operation computes the top-left corners of all $N \times M$ intersection rectangles. Similarly for bottom-right corners. The intersection area is computed with clamping to handle non-overlapping boxes (intersection area = 0). Finally, IoU is intersection divided by union, computed in a fully vectorized way.

This efficient computation is crucial because we might compute IoU between 1,764 anchors and multiple ground truth boxes, potentially thousands of comparisons, and we do this every training iteration.

Assigning Anchors to Ground Truth

The `assign_anchors` method determines labels for all anchors. It initializes all labels to -1 (ignore), then computes IoU between all anchors and all ground truth boxes. For each anchor, we find the maximum IoU across all GT boxes and the index of that best-matching

box. The labeling rules are straightforward: if $\text{max IoU} < 0.3$, it's background (label 0); if $\text{max IoU} \geq 0.7$, it's foreground (label 1); everything in between is ignored (stays at -1). This creates a clean separation with an ambiguous middle region that doesn't contribute to training.

For positive anchors, we need regression targets. We extract the matched ground truth boxes using the `max_idx` indices and call `compute_bbox_targets` to compute the transformation from anchor to ground truth. This is the same parameterization as earlier R-CNN variants: normalized center offsets and log-space scale changes.

The Training Loss Computation

The `__call__` method computes the actual loss for a batch. It receives predictions (objectness and `bbox_deltas`), anchors, and a list of ground truth boxes (one tensor per image in the batch). First, we flatten the spatial dimensions of the predictions - objectness becomes $[B, N]$ where $N = H \times W \times \text{num_anchors}$, and `bbox_deltas` becomes $[B, N, 4]$.

For each image in the batch, we assign anchors to ground truth, getting labels and `bbox_targets`. Then comes a crucial step: mini-batch sampling. We can't use all anchors because there are too many and they're heavily imbalanced (most are background). We randomly sample up to 128 positive anchors and up to 128 negative anchors (or fewer if not enough exist). Ignored anchors (label -1) never contribute to training.

The classification loss is binary cross-entropy with logits, comparing predicted objectness scores against binary labels (1 for object, 0 for background) for the sampled anchors. The regression loss is smooth L1, but only for positive anchors - background anchors have no associated box to regress to. We average losses across the batch and return both components separately so we can monitor them during training.

This loss function embodies the multi-task nature of RPN: simultaneously learn to classify anchors (object vs. background) and refine them (predict box adjustments). The sampling strategy ensures balanced training despite the severe class imbalance inherent in dense anchor placement.

The FasterRCNN Class: The Complete Pipeline

The `FasterRCNN` class assembles all components into a complete end-to-end detection system. It combines a CNN backbone, anchor generation, and the RPN into a single trainable model.

The initialization builds the network architecture. The backbone here is simplified for demonstration - just a sequence of convolutions that downsample the input and increase channels. In practice, you'd use ResNet or another modern architecture. The important thing is that it outputs feature maps with a known channel dimension (512 here) and a known stride (16 - the input is downsampled by a factor of 16).

We instantiate the anchor generator with the standard configuration (3 scales, 3 ratios, stride 16), giving us 9 anchors per location. The RPN gets initialized with the backbone's output

channels and the number of anchors. This modular design is clean: the backbone extracts features, the anchor generator defines the search space, and the RPN makes predictions within that space.

Forward Pass: Training and Inference Modes

The `forward` method handles both training and inference. It extracts the batch size and image dimensions, then runs the images through the backbone to get shared CNN features - this is the key efficiency gain over R-CNN. We generate anchors based on the feature map size and original image size. Then the RPN makes its predictions: objectness scores and `bbox_deltas` for all anchors.

The method branches based on mode. During training (when `gt_boxes` is provided), we compute the RPN loss using our loss function and return a dictionary of losses. During inference, we decode the RPN predictions into actual proposal boxes and return them. This dual-mode behavior is standard in PyTorch detection models - the same forward pass with different behavior based on whether we're training or testing.

Decoding Proposals from RPN Predictions

The `decode_proposals` method converts RPN outputs into concrete bounding box proposals. It flattens the spatial dimensions to get $[B, N]$ objectness scores and $[B, N, 4]$ `bbox_deltas`. For each image in the batch, we apply sigmoid to objectness scores to get probabilities, then select the top-k anchors by score. Typically $k=2000$, matching the number of proposals Selective Search would generate.

For these top-k anchors, we extract their `bbox_deltas` and apply them to the anchor coordinates to get refined boxes. The `apply_deltas` method implements the inverse of the `bbox` regression parameterization: we convert anchors to center coordinates and dimensions, apply the predicted deltas (adding the tx and ty offsets scaled by box dimensions, exponentiating the tw and th log-scale factors), then convert back to corner coordinates.

The result is a list of proposal boxes for each image, ready to be fed into a detection head (RoI pooling + classification/regression). These proposals are learned - the network discovered them by examining the image features, not through hand-crafted algorithms.

The Demo Function: Seeing Everything Together

The `demo_faster_rcnn` function provides a complete demonstration of the full system. It creates a FasterRCNN model and dummy training data (images and ground truth boxes). The demonstration proceeds in clear stages, each with explanatory output.

First, it prints the model architecture summary, establishing what components are involved. Then it trains the RPN for 5 epochs, using SGD with momentum. For each epoch, it does a forward pass to get losses, computes the total loss as the sum of classification and

regression losses, and performs backpropagation. The printed losses show the network learning to classify anchors and refine them.

After training, the model switches to evaluation mode for inference. It generates proposals for the test images and reports how many proposals were generated. The key teaching moment comes in the evolution comparison, showing the progression from Selective Search (hand-crafted) in R-CNN and Fast R-CNN to the learned RPN in Faster R-CNN. The computational comparison is striking: Selective Search takes 2 seconds per image, RPN takes 0.01 seconds - a $200\times$ speedup, plus the proposals are learned and thus better suited to the detection task.

The Fundamental Innovation: Learning the Search Space

Understanding the anchor generation and RPN classes reveals Faster R-CNN's crucial innovation. R-CNN and Fast R-CNN outsourced the "where to look" question to Selective Search, a separate, untrainable algorithm. Faster R-CNN internalizes this question: the RPN learns where objects are likely to be by examining learned features.

The anchor mechanism makes this feasible. Instead of asking the network to predict arbitrary boxes (an unconstrained, difficult problem), we give it templates and ask it to answer two simpler questions: "Which templates look like objects?" and "How should we adjust them?" This constraint dramatically simplifies learning. The anchors provide a discrete, structured search space; the RPN learns to navigate it.

The multi-scale, multi-aspect-ratio design ensures comprehensive coverage. Small objects will match small-scale anchors, large objects will match large-scale ones. Tall objects (like pedestrians) will match high-aspect-ratio anchors, while wide objects (like cars) will match low-aspect-ratio anchors. The RPN doesn't need to know about object categories - it just learns the general concept of "objectness" across all scales and shapes.

This architecture became foundational. The anchor concept appears in RetinaNet, SSD, and YOLO (though YOLO uses anchor-like priors). Even anchor-free methods like FCOS and DETR are partly reactions to anchor-based detection, showing how central this idea became. Understanding these classes deeply - how anchors are generated, how RPN makes predictions, how the loss function trains it - provides the foundation for understanding all modern object detection. That's why we start here: master anchors and RPN, and you've mastered the core paradigm of detection.

A Detailed Class-by-Class Walkthrough of FPN and RetinaNet

Let me guide you through Feature Pyramid Networks and RetinaNet, two pivotal innovations that addressed fundamental limitations in object detection. FPN solved the multi-scale detection problem elegantly, while RetinaNet proved that one-stage detectors could finally match two-stage accuracy by solving the class imbalance problem. Together, they represent a major leap forward in detection architecture.

The Historical Context: Why We Needed FPN

Before diving into the code, understand the problem FPN solved. In CNNs, there's an inherent tension: early layers have high spatial resolution but weak semantic information (they see edges and textures but don't understand "what" things are), while deep layers have strong semantics but low resolution (they understand objects but have lost spatial detail). For detection, this creates a dilemma. If we use only deep features (C5 from ResNet, typically 7×7), we get strong semantics but struggle with small objects that become tiny or invisible at that resolution. If we use shallow features (C3, typically 28×28), we have resolution but the features don't understand objects well enough for reliable detection.

Previous approaches used image pyramids - literally resizing the image to multiple scales and running the detector on each. This works but is computationally expensive. FPN's insight was profound: we can build a feature pyramid inside the network itself, getting multi-scale features from a single forward pass. The key innovation is the top-down pathway that brings semantic information from deep layers back to shallow ones, creating features that have both high resolution and strong semantics.

The SimpleFPN Class: Multi-Scale Feature Fusion

The `SimpleFPN` class implements the core FPN architecture in a clear, educational way. It demonstrates how we transform the backbone's multi-scale features (C3, C4, C5) into a feature pyramid (P3, P4, P5) where every level has both good resolution for its scale and strong semantic understanding.

The initialization defines the architecture's two key components. First, the lateral connections - three 1×1 convolutions that reduce the channel dimensions of C3, C4, and C5 to a unified output dimension (typically 256). These 1×1 convs serve a crucial purpose: they align the channel dimensions across different backbone stages so we can add features together. C3 might have 512 channels, C4 has 1024, and C5 has 2048, but we need them all to have the same number of channels (256) to combine them. The 1×1 conv also acts as a learned projection that selects which information from each stage is most useful for detection.

Second, we have output 3×3 convolutions for each pyramid level. These "smoothing" convolutions reduce aliasing artifacts that can occur from the upsampling operations in the top-down pathway. When you upsample using nearest neighbor interpolation, you can get

blocky features; the 3×3 conv smooths these out while also providing additional representational capacity at each level.

The Forward Pass: Building the Pyramid

The `forward` method implements the FPN's three-stage construction with detailed logging to aid understanding. It accepts features either as a dictionary (with keys 'C3', 'C4', 'C5') or as a list - this flexibility makes the code more practical. The logging output is pedagogically valuable, showing the shapes at each stage and annotating what each represents (stride, resolution characteristics, semantic strength).

Stage one applies the lateral connections. We pass each backbone feature through its corresponding 1×1 conv, producing lat_3 , lat_4 , and lat_5 , all with 256 channels but preserving their original spatial dimensions. This alignment is the foundation that makes the top-down pathway possible.

Stage two builds the top-down pathway, and this is where FPN's magic happens. We start from the deepest level P_5 , which is just lat_5 - it already has strong semantics from being deep in the network. To build P_4 , we upsample P_5 to match lat_4 's spatial dimensions using nearest neighbor interpolation, then add it element-wise to lat_4 . This addition is the fusion operation: lat_4 contributes its spatial precision (it came from a higher-resolution stage), while upsampled P_5 contributes its semantic understanding. The result, P_4 , has medium resolution and strong semantics.

We repeat this for P_3 : upsample P_4 to match lat_3 's spatial dimensions and add. Now P_3 has the original high resolution from C_3 (28×28 for a 224×224 input) plus the semantic information that has flowed down from C_5 through P_5 and P_4 . This is the key insight - by repeatedly upsampling and adding, we create a pathway for semantic information to flow from deep to shallow, enriching the shallow features without losing their spatial precision.

Stage three applies the smoothing convolutions. Each pyramid level passes through its 3×3 conv, producing the final outputs P_3 , P_4 , and P_5 . These convolutions serve multiple purposes: they smooth the upsampling artifacts, they provide additional non-linear transformation capacity, and they give each level a chance to refine its fused representation. The logging explicitly marks these as having both high resolution (for their respective scales) and strong semantics (from the top-down pathway).

The returned list $[P_3, P_4, P_5]$ represents a proper feature pyramid: P_3 for detecting small objects (stride 8, roughly 8-64 pixel objects), P_4 for medium objects (stride 16, 64-128 pixels), and P_5 for large objects (stride 32, 128+ pixels). Each level has 256 channels and semantic understanding comparable to the deepest layer.

The SimpleBackbone Class: Generating Multi-Scale Features

The `SimpleBackbone` class provides a simplified CNN that generates the multi-scale features FPN needs. In practice, you'd use ResNet or another modern architecture, but this simplified version clearly shows the structure needed for FPN.

The initialization defines a sequence of convolutional layers with progressively increasing stride and channels. The first conv with stride 2 downsamples to half resolution, the max pool with stride 2 downsamples again (total stride 4 so far). Layer2 maintains this stride but increases channels to 512, producing C3 with stride 8. Layer3 adds another max pool (stride 16 total) and increases channels to 1024 for C4. Layer4 gives us stride 32 and 2048 channels for C5.

This architecture mirrors ResNet's structure: progressively deeper stages with increasing stride and channels. The important property is that we extract intermediate features at multiple scales. C3, C4, and C5 represent the same image at different resolutions and with different levels of semantic abstraction.

The `forward` method processes the input through this backbone and returns the multi-scale features as a dictionary. This clean interface makes it easy for FPN to consume the features - it just indexes into the dict to get each scale. The explicit extraction and return of intermediate features is what enables FPN; earlier architectures might only return the final output, but FPN requires access to these intermediate representations.

The FPN Architecture Visualization

The `visualize_fpn_architecture` function creates a comprehensive diagram showing FPN's complete data flow. This visualization is crucial for teaching because FPN's architecture is more complex than previous detectors - it has multiple pathways and connections that need to be understood spatially.

The function sets up a matplotlib canvas and draws FPN's components as connected boxes with arrows. The bottom-up pathway (backbone) is drawn on the left, showing C3, C4, and C5 as progressively smaller boxes (representing their decreasing spatial resolution) with increasing channels. Each box is colored differently and labeled with its properties: channel count, stride, and scale. The arrows between backbone levels show the forward progression through the network.

The lateral connections are drawn as small yellow boxes (representing 1×1 convs) with arrows from the backbone features. These visually connect the bottom-up and top-down pathways, showing how information flows horizontally across the network at each scale.

The top-down pathway is drawn with purple boxes and importantly includes dashed red arrows showing the upsampling operations. These arrows point upward (from deeper to shallower), visually reinforcing that information flows top-down. The annotations next to these arrows say "upsample," making it clear that this isn't just a connection but a resolution-increasing operation.

The output convolutions are shown as gray boxes (representing 3×3 convs), and finally, the pyramid features P3, P4, and P5 are shown as orange boxes, prominently labeled with their

unified 256-channel dimension. Arrows connect these to detection heads (shown on the right), illustrating that each pyramid level feeds into its own detection head for multi-scale detection.

The legend at the bottom summarizes the key benefits: combining high resolution with strong semantics, multi-scale detection without pyramids, and channel unification. This visualization serves as a reference diagram for the entire FPN concept, showing in a single image what pages of text struggle to convey.

The Demonstration Functions: Teaching FPN

The `demonstrate_fpn_features` function provides a complete demonstration of FPN's operation with detailed output. It creates the backbone and FPN, processes a dummy input, and prints extensive information about what happens at each stage. The output is structured to teach: first showing the input image, then the backbone features (with annotations about their resolution and semantic strength), then the FPN forward pass with detailed logging of each operation, and finally the output pyramid with its unified properties.

The key insight is explicitly stated in the output: all pyramid levels have 256 channels (unified), but P3 has both high resolution (28×28) and strong semantics (brought down from C5). This breaks the traditional CNN trade-off. The demonstration also maps each pyramid level to what object sizes it should detect, helping students understand how multi-scale detection works in practice.

The `compare_with_without_fpn` function provides the crucial context for why FPN matters. It contrasts single-scale detection (using only C5) with multi-scale FPN detection. The single-scale approach struggles because small objects disappear in the 7×7 resolution of C5. FPN solves this by using P3's 28×28 resolution for small objects while ensuring P3 still has strong semantic understanding via the top-down pathway. The performance numbers make this concrete: +8% AP on small objects is a massive improvement.

The Historical Significance of FPN

Understanding FPN's implementation reveals why it was transformative. Before FPN, multi-scale detection required either image pyramids (slow) or using features from a single scale (inaccurate for objects of varying sizes). FPN proved that you could build a feature pyramid inside the network using features you're computing anyway, getting multi-scale capability essentially for free. The top-down pathway was the key innovation - it's not just using multi-scale features (people had tried that), but enriching those features with semantic information from deeper layers.

FPN became foundational immediately. Mask R-CNN built on it, RetinaNet used it, and virtually all modern detectors incorporate FPN or FPN-like multi-scale feature fusion. The pattern of lateral connections plus top-down pathways appears everywhere in computer vision now, extending beyond detection to segmentation, pose estimation, and other tasks. Understanding this code deeply means understanding a core design pattern in modern vision architectures.

The FocalLoss Class: Solving Class Imbalance

Now we shift to RetinaNet and the focal loss, which solved a different but equally important problem. The `FocalLoss` class implements the loss function that made one-stage detectors viable for high-accuracy detection.

The initialization stores two hyperparameters that define the focal loss behavior. Alpha (default 0.25) is a balancing factor for positive vs negative examples, similar to weighted cross-entropy. Gamma (default 2.0) is the focusing parameter that controls how much we down-weight easy examples. These aren't arbitrary values - the RetinaNet paper extensively ablated different settings and found $\alpha=0.25$ and $\gamma=2.0$ worked best across datasets.

The Forward Pass: Computing Focal Loss

The `forward` method implements the focal loss formula, and understanding each step reveals why it works. First, we apply sigmoid to get probabilities from logits. This gives us p , the predicted probability for each class. For binary classification per class (which is how we handle multi-class in detection), we compute p_t , the probability of the correct class. If the target is 1 (positive), p_t is just p . If the target is 0 (negative), p_t is $1-p$. This ensures p_t is always the probability of being correct.

The focal loss then modifies standard cross-entropy with two factors. The focal weight $(1 - p_t)^\gamma$ is the key innovation. For easy examples where p_t approaches 1 (the model is confident and correct), this term approaches 0, dramatically reducing the loss contribution. For hard examples where p_t is low (the model is uncertain or wrong), this term approaches 1, preserving the full loss. The gamma parameter controls the strength of this effect: $\gamma=0$ gives standard cross-entropy (no focusing), $\gamma=2$ provides strong focusing (easy examples at $p_t=0.9$ get their loss reduced by 99%).

The alpha weighting provides additional balance. We multiply by α for positive examples and $(1-\alpha)$ for negatives. With $\alpha=0.25$, positive examples get 25% weighting and negatives get 75%. This might seem backwards - shouldn't we weight rare positives more? But remember, there are 1000 \times more negatives, so the $(1-\alpha)$ factor still gives negatives much higher total weight. The α factor is actually compensating for the class imbalance, not causing it.

The final loss is the product of `alpha_t`, `focal_weight`, and cross-entropy. We average across all examples to get a scalar loss. This loss function is drop-in compatible with cross-entropy - it takes the same inputs and returns a scalar - but it has radically different behavior on easy vs hard examples.

Visualizing Focal Loss: Understanding the Effect

The `visualize_focal_loss` function creates three panels that together explain how focal loss works and why it's necessary. This visualization is pedagogically crucial because focal loss is mathematically simple but conceptually subtle.

Panel 1 plots loss curves for cross-entropy and focal loss with various gamma values. The x-axis is p , the probability of the correct class, and the y-axis is the loss. For cross-entropy ($\gamma=0$), even at $p=0.9$ (a confident, correct prediction), the loss is still $-\log(0.9) \approx 0.1$. These "easy" examples still contribute to the loss. For focal loss with $\gamma=2$, at $p=0.9$, the loss is $(1-0.9)^2 \times (-\log(0.9)) \approx 0.001$ - a 100x reduction! The plot shows this dramatic difference, with higher gamma values creating even steeper curves that approach zero faster for confident predictions.

The highlighting on this panel identifies "easy" ($p > 0.8$) and "hard" ($p < 0.5$) regions, making it visually clear that focal loss heavily down-weights the easy region while preserving loss in the hard region. Students can literally see how the loss behaves differently in different confidence regimes.

Panel 2 zooms in on the modulating factor $(1-p)^\gamma$ for a specific easy example ($p=0.9$). It plots how this factor changes with gamma. At $\gamma=0$, the factor is 1 (no modulation). At $\gamma=2$, the factor is 0.01 (99% reduction). At $\gamma=5$, it's 0.00001 (99.999% reduction). The annotation points out that $\gamma=2$ reduces the weight to 1%, which is the sweet spot - aggressive enough to down-weight easy examples but not so aggressive that it completely ignores them.

Panel 3 is perhaps the most important for understanding the practical impact. It simulates a realistic class distribution: 99.9% easy negatives (the model correctly identifies them as background with high confidence) and 0.1% hard positives (the model struggles to detect actual objects). For each group, we compute the total loss contribution under cross-entropy and focal loss. The bar chart shows a dramatic reversal: with cross-entropy, easy negatives contribute 836 units of loss while hard positives contribute only 1 unit - the easy examples dominate training by 836:1! With focal loss, easy negatives contribute only 0.3 units while hard positives contribute 1.2 units - now hard examples dominate 4:1.

The annotations make this explicit: "CE: Easy negatives dominate training!" vs "FL: Hard examples get more focus!" This panel transforms the abstract mathematical formula into a concrete practical insight: focal loss fundamentally changes what the network pays attention to during training.

The SimpleRetinaNet Class: One-Stage Detection with Focal Loss

The [SimpleRetinaNet](#) class brings together FPN's multi-scale features and focal loss into a complete one-stage detector. The architecture is elegantly simple, which is part of RetinaNet's appeal - it proves you don't need complex machinery, just the right loss function.

The initialization sets up the detection subnets. RetinaNet uses a crucial design choice: the classification and box regression heads are separate subnets, but they're shared across all FPN levels. This sharing is efficient (fewer parameters) and provides a useful inductive bias (the same patterns appear at all scales).

The classification subnet is built as a sequence of four 3×3 convolutions with ReLU, followed by a final 3×3 conv that produces the predictions. Each intermediate conv has 256 channels

(matching FPN's output channels), providing substantial representational capacity. The final conv outputs $\text{num_anchors} \times \text{num_classes}$ channels - for each anchor at each spatial location, we predict a score for each class. The depth (four convs) is a deliberate choice - deeper than just a single conv head but not so deep as to be slow or hard to train.

The box subnet has an identical structure but outputs $\text{num_anchors} \times 4$ channels (four bbox coordinates per anchor). This symmetry is aesthetically pleasing and practically effective - both tasks (classification and regression) benefit from the same depth of processing.

Weight Initialization: Critical for Focal Loss

The `_init_weights` method includes a crucial detail that's easy to overlook but vital for training with focal loss. All convolutional layers get standard initialization: Gaussian weights with $\text{std}=0.01$ and zero biases. But the final classification layer gets special treatment: its bias is initialized to $-\log((1-\pi)/\pi)$ where $\pi=0.01$ is the prior probability of an object.

Why this specific initialization? At the start of training, before the network has learned anything, it will output roughly zero for all logits (due to small random weights). Applying sigmoid to zero gives 0.5 - the network predicts 50% probability for every class. But with 1000:1 class imbalance (0.1% positives), this is wildly wrong and causes massive loss. By setting the bias to $-\log(99)$, the initial sigmoid output becomes $1/(1+99) \approx 0.01$, matching the actual class frequency. This initialization is critical for focal loss - it ensures the network starts from a reasonable place rather than having to overcome a huge initial loss from imbalanced predictions.

Forward Pass: Predictions Across Scales

The `forward` method processes FPN features to generate predictions at all scales. It receives a list of feature maps [P3, P4, P5, ...], one per pyramid level. For each feature map, it applies both the classification subnet and the box subnet. The subnets are the same `nn.Module` objects applied to different inputs - this is the weight sharing across scales.

The outputs are lists of tensors, one per FPN level. Each classification tensor has shape [B, $\text{num_anchors} \times \text{num_classes}$, H, W] where H and W are the spatial dimensions of that FPN level. Each box tensor has shape [B, $\text{num_anchors} \times 4$, H, W]. These lists preserve the multi-scale structure - we know which predictions came from which scale, which matters for both training (assigning ground truth to the right scale) and inference (interpreting detection sizes).

This multi-scale prediction is RetinaNet's strength. Unlike Faster R-CNN, which uses RPN to filter proposals before detection, RetinaNet makes dense predictions at every spatial location across all scales. With 9 anchors per location and 5 FPN levels, that might be 100,000+ predictions per image. This density is only tractable because focal loss handles the massive class imbalance that results from such dense prediction.

The Focal Loss Demonstration: Making It Concrete

The `demo_focal_loss_effect` function demonstrates focal loss's impact with realistic numbers. It sets up the scale of the problem: 100,000 anchor boxes per image with only 100 positives - a 1000:1 imbalance that's typical in detection.

The simulation creates two groups of predictions. Easy negatives are modeled with logits drawn from a distribution shifted positive (mean +3), representing high confidence that these are background. Hard positives are modeled with logits shifted negative (mean -1), representing uncertainty about whether these contain objects. This mimics early training, when the network struggles with objects but is confident about obvious background regions.

Computing cross-entropy loss for each group reveals the problem. Easy negatives, despite being correctly classified, still contribute substantial loss because cross-entropy doesn't have a "good enough" threshold - even 99% confidence produces non-zero loss. With 99,900 easy negatives, their total loss overwhelms the 100 hard positives by 836:1. The network spends most of its gradient budget improving already-good predictions rather than fixing mistakes.

With focal loss ($\gamma=2$, $\alpha=0.25$), the balance shifts dramatically. The $(1-p)^2$ modulating factor reduces easy negative loss by ~99%, while hard positive loss is preserved. The printed output shows the reversal: easy negatives now contribute less than hard positives. The network can focus on the hard examples that actually need improvement.

The explanatory text breaks down the focal loss formula into its components, explaining what each term does. The modulating factor handles easy/hard imbalance through its exponential dependence on confidence. Alpha handles class imbalance through simple weighting. Together, they make training viable despite severe imbalance on both dimensions (class distribution and example difficulty).

The Complete RetinaNet Demonstration

The `demo_retinanet_complete` function orchestrates a comprehensive demonstration that builds understanding progressively. First, it visualizes the focal loss curves, giving students the mathematical intuition. Second, it demonstrates the practical effect with simulated data, showing the loss contribution reversal. Third, it explains RetinaNet's architecture and shows how the pieces fit together.

The architecture summary is pedagogically structured: backbone (ResNet) → FPN (multi-scale features) → subnets (shared heads) → focal loss (handle imbalance). Each component is introduced in dependency order. The forward pass demonstration shows the actual tensor shapes at each FPN level, making the multi-scale prediction concrete.

The performance comparison drives home RetinaNet's significance: 39.1% AP, beating Faster R-CNN's 36.2%. This is historically important - RetinaNet was the first one-stage detector to match two-stage accuracy. Previous one-stage detectors (like early YOLO or SSD) were faster but less accurate. RetinaNet proved that with the right loss function, one-stage could be both fast and accurate.

The key takeaways explicitly state the insights: class imbalance is the problem, focal loss is the solution, FPN provides multi-scale features, and the combination achieves state-of-the-art results. This explicit summary helps students synthesize the technical details into high-level understanding.

The Broader Impact: FPN and RetinaNet's Legacy

Understanding these implementations in detail reveals why FPN and RetinaNet were transformative. FPN solved multi-scale detection elegantly - the top-down pathway is a simple idea but it works beautifully, giving us high-resolution features with strong semantics. Every modern detector uses FPN or a variant because it's simply the right way to handle scale variation.

RetinaNet's focal loss solved the class imbalance problem that plagued one-stage detectors. By down-weighting easy examples, it made dense prediction viable. The loss function itself is remarkably simple - just add a modulating factor to cross-entropy - but the impact was profound. Suddenly, one-stage detectors could match two-stage accuracy without sacrificing speed.

Together, these innovations enabled the current generation of detectors. FCOS builds on FPN and focal loss but removes anchors entirely. EfficientDet uses FPN's idea but makes it bidirectional. DETR uses transformers but still needs multi-scale features. Modern YOLO versions incorporate focal loss variants. The patterns established by FPN (multi-scale feature fusion) and RetinaNet (hard example mining through loss design) appear everywhere.

That's why we study these implementations carefully. FPN teaches us that architectural innovations can come from simple ideas executed well - just upsample and add, but do it in the right way. RetinaNet teaches us that loss function design is as important as architecture - the same model with different losses has radically different behavior. These lessons apply far beyond object detection, influencing how we think about multi-scale problems and class imbalance across computer vision and machine learning generally.

A Detailed Class-by-Class Walkthrough of YOLO

Let me guide you through the YOLO (You Only Look Once) implementation, which represents a radical paradigm shift in object detection. While R-CNN, Fast R-CNN, and Faster R-CNN all follow a region-based approach (propose regions, then classify them),

YOLO asked a fundamentally different question: what if we treat detection as a single regression problem, directly predicting bounding boxes and class probabilities from image pixels in one shot?

The Revolutionary Paradigm: Detection as Regression

Before diving into the code, understand why YOLO was revolutionary. Previous detectors followed a two-step process: first generate proposals (whether through Selective Search or RPN), then classify and refine them. This separation makes intuitive sense - first find where objects might be, then figure out what they are. But it's also inefficient and indirect. YOLO's insight was profound: a single neural network can look at the entire image once and directly predict all bounding boxes and their classes simultaneously. No proposals, no region pooling, no multi-stage pipeline - just one forward pass.

The key to making this work is the grid-based formulation. YOLO divides the image into an $S \times S$ grid (typically 7×7). Each grid cell is responsible for detecting objects whose center falls within that cell. Each cell predicts B bounding boxes (typically 2) and confidence scores for those boxes, plus C class probabilities. The output is a tensor of shape $[S, S, B \times 5 + C]$, encoding all detections in one structured prediction. This formulation is elegant because it naturally handles multiple objects (different grid cells detect different objects) while maintaining spatial coherence (nearby predictions come from nearby cells).

The SimpleBackbone Class: Feature Extraction for Detection

The `SimpleBackbone` class implements a convolutional neural network that extracts features for YOLO's detection head. While YOLO v1 used a custom architecture inspired by GoogLeNet, this simplified version demonstrates the key principles: progressively downsample the image while increasing feature channels, ultimately producing a feature map that preserves enough spatial resolution for localization while having sufficient semantic understanding for classification.

The initialization builds a sequential network with a clear downsampling pattern. The first convolutional block uses a 7×7 kernel with stride 2, immediately reducing the 224×224 input to 112×112 . This large receptive field in the first layer helps the network capture broader context from the start. BatchNorm and LeakyReLU follow - LeakyReLU with slope 0.1 is YOLO's standard choice, providing non-linearity while avoiding the dying ReLU problem. A max pool with stride 2 further reduces to 56×56 .

The subsequent blocks follow a pattern: convolution to increase channels, normalization, activation, sometimes additional convolutions at the same resolution, then max pooling to downsample. This creates a feature hierarchy where each stage operates at half the resolution of the previous one but with double the channels. The channel progression ($64 \rightarrow 128 \rightarrow 256 \rightarrow 512$) reflects increasing abstraction - early layers capture low-level features, later layers understand high-level semantics.

The final output is 512 channels at 14×14 resolution (for 224×224 input). This relatively high resolution compared to classification networks is deliberate - YOLO needs spatial precision for localization. A 7×7 feature map (YOLO's typical grid size) would lose too much spatial information, making it hard to precisely localize objects. The 14×14 features are then adaptively pooled to 7×7 in the detection head.

The `forward` method is straightforward - it just passes the input through the sequential network and returns the features. The simplicity belies the importance: this backbone must learn features that support both localization (where is the object?) and classification (what is it?) simultaneously.

The YOLODetector Class: Unified Detection

The `YOLODetector` class is where YOLO's unified detection happens. It combines the backbone with a detection head that predicts all bounding boxes and classes in a single forward pass. Understanding this class reveals how YOLO achieves its speed and simplicity.

The initialization stores the fundamental YOLO parameters: number of classes, grid size S, and number of boxes per cell B. These define the output structure. The backbone is instantiated with the feature dimension noted (512 channels). Then comes the detection head, which is beautifully simple: a 3×3 conv to increase capacity ($512 \rightarrow 1024$ channels), followed by dropout for regularization, and finally a 1×1 conv to produce predictions.

The output channel count calculation is the key: `num_boxes * 5 + num_classes`. For each of the B boxes, we predict 5 values: x, y, w, h (bounding box coordinates) and confidence (objectness score). Additionally, for the cell (not per box), we predict C class probabilities. With B=2 boxes and C=3 classes, we get $2 \times 5 + 3 = 13$ output channels. This compact representation is what makes YOLO efficient.

The adaptive pooling layer explicitly sets the spatial dimensions to the grid size (7×7). This is crucial because the backbone might output features at a different resolution depending on input size. Adaptive pooling ensures we always get exactly $S \times S$ spatial dimensions, matching YOLO's grid structure.

Weight initialization uses Kaiming initialization for the detection head, appropriate for networks with LeakyReLU activations. This ensures the network starts with reasonable gradient magnitudes, facilitating training.

The Forward Pass: Grid-Based Prediction

The `forward` method implements YOLO's core computation. It takes images [B, 3, H, W] and produces structured predictions [B, S, S, B, 5+C]. Understanding this transformation reveals YOLO's elegance.

First, features are extracted via the backbone, producing [B, 512, ~14, ~14] features. These encode the image at a resolution that balances spatial precision with semantic

understanding. The adaptive pooling then reshapes to exactly [B, 512, S, S] - forcing the features to match our grid structure.

The detection head processes these features, outputting [B, B×5+C, S, S]. This tensor is then permuted to [B, S, S, B×5+C] - moving the spatial dimensions before the feature dimension. This reordering makes it easier to think about: "for each batch item, for each grid cell (S×S), we have B×5+C predictions."

The crucial reshaping comes next: [B, S, S, B×5+C] → [B, S, S, B, 5+C]. We're separating the B×5+C predictions into B boxes, each with 5+C values. Now the structure is explicit: batch, grid_y, grid_x, box_index, prediction_values.

The activation functions applied at the end are critical for YOLO's formulation. The x, y coordinates get sigmoid, constraining them to [0, 1] - they represent offsets within the cell, not absolute positions. The confidence also gets sigmoid, giving a probability-like score. The class probabilities get softmax, ensuring they're a proper probability distribution over classes. Width and height are left as-is (they'll be used as absolute values relative to the image size).

This output structure is YOLO's signature: a dense tensor encoding all predictions simultaneously, organized by spatial location and box index. There's no filtering, no proposals, no selection - just a complete prediction for every position.

The YOLOLoss Class: Multi-Task Training

The `YOLOLoss` class implements YOLO's multi-component loss function, which is more complex than it first appears. The challenge is balancing multiple objectives: localization accuracy, confidence scores, and classification, while handling the fact that most grid cells contain no objects.

The initialization stores the grid and box parameters, plus two crucial weighting factors. `Lambda_coord` (default 5.0) upweights the localization loss - this compensates for the fact that localization errors affect fewer cells (only those with objects) compared to confidence errors (all cells). `Lambda_noobj` (default 0.5) downweights the confidence loss for cells without objects - without this, the overwhelming number of no-object cells would dominate training, causing the network to predict low confidence everywhere.

Computing IoU for Responsibility

The `compute_iou` method calculates intersection over union between predicted and ground truth boxes. YOLO uses IoU to determine which of the B predicted boxes in a cell is "responsible" for detecting an object - it's whichever predicted box has highest IoU with the ground truth.

The implementation converts boxes from [x, y, w, h] (center coordinates and dimensions) to [x1, y1, x2, y2] (corners). For YOLO, x and y are cell-relative offsets, but w and h are absolute dimensions (relative to image size), so the conversion uses half-widths and half-heights around the center. Computing intersection area uses the standard max/min approach for finding overlapping rectangles, with clamping to handle non-overlapping boxes.

Union is sum of areas minus intersection. The IoU calculation includes epsilon (1e-6) for numerical stability when boxes don't overlap.

The Loss Computation: Five Components

The `forward` method computes YOLO's loss as a sum of five terms, each addressing a specific aspect of detection. Understanding each component reveals YOLO's training strategy.

First, we extract components from predictions and targets. Both have the same shape [B, S, S, B, 5+C], making the extraction parallel. The responsibility mask is determined by `target_conf > 0` - only boxes that were assigned an object (responsibility was given during target encoding) contribute to certain losses. This mask is critical because it implements YOLO's responsibility concept: only the designated box per object should be trained.

Localization Loss (xy): For the x, y coordinates, we use squared error between predicted and target, weighted by `lambda_coord` (5.0) and only for responsible boxes (`obj_mask`). This loss trains the network to predict accurate box centers. The `lambda_coord` upweighting ensures localization gets sufficient gradient signal despite affecting fewer cells than other losses.

Localization Loss (wh): For width and height, YOLO uses a clever formulation: instead of directly regressing w and h, we regress their square roots. Why? Because errors in large boxes should be penalized less than errors in small boxes. A 10-pixel error matters much more for a 20-pixel box than a 200-pixel box. Using square root makes the loss more scale-invariant: $\sqrt{(\text{large})} - \sqrt{(\text{small})}$ differences are more similar in magnitude than $(\text{large}) - (\text{small})$ differences.

Confidence Loss (object present): When an object is present, we want predicted confidence to match the actual IoU between predicted and ground truth boxes (in this simplified implementation, we just use 1.0 as the target confidence). The squared error trains the network to be confident when it's responsible for an object. Only responsible boxes contribute to this loss.

Confidence Loss (no object): When no object is present, we want predicted confidence to be 0. This loss is weighted by `lambda_noobj` (0.5) to prevent the overwhelming majority of no-object cells from dominating training. Without this downweighting, the network would simply predict 0 confidence everywhere to minimize loss. The reduced weight allows object detections to provide sufficient gradient signal.

Classification Loss: For cells containing objects, we predict class probabilities and compare to ground truth using squared error (YOLO v1 used squared error rather than cross-entropy). The `cell_obj_mask` ensures we only compute classification loss for cells that actually contain objects - there's no meaningful class prediction for empty cells.

The total loss sums these five components and normalizes by batch size. The return includes both the total loss (for backpropagation) and a dictionary of individual components (for monitoring training). This decomposition is pedagogically valuable - it shows exactly what aspects of detection the network is learning.

The SyntheticObjectDataset Class: Generating Training Data

The `SyntheticObjectDataset` class generates synthetic images with geometric shapes, providing a controlled environment for understanding YOLO's behavior without the complexity of real datasets. This is an excellent pedagogical choice because students can see exactly what the network is trying to detect.

The initialization defines the dataset parameters: number of samples, image size, grid size (matching YOLO's grid), number of boxes and classes, and maximum objects per image. It also sets up class names (circle, square, triangle) and distinctive colors for each class. These shapes are chosen because they're easy to generate, visually distinct, and suitable for demonstrating multi-class detection.

Drawing Shapes: Creating Visual Diversity

The `draw_shape` method implements shape rendering using NumPy operations. Each shape type requires different rasterization logic. For circles, we create coordinate grids and use the distance formula - points within radius/2 of the center are marked. For squares, we simply fill a rectangular region. For triangles, we use matplotlib's Path class to check point containment - this is more complex but demonstrates how to rasterize arbitrary polygons.

The color assignment uses the class-specific colors, making it visually obvious which class each shape belongs to. This visual distinctiveness helps when debugging - you can immediately see if the network is confusing classes.

Getting Items: Image and Target Generation

The `__getitem__` method creates a training sample. It starts with a blank off-white image (0.95 intensity, not pure white, for better contrast), then randomly generates 1 to `max_objects` shapes. For each object, we randomly sample its class, size (30-80 pixels), and position (ensuring it fits within the image). The `draw_shape` method renders it, and we record its normalized coordinates (`x`, `y`, `w`, `h` all in [0, 1] range) and class.

The image is converted to a tensor [3, H, W] with channel-first format. The objects list is encoded into YOLO target format using `encode_target`. Both are returned along with the raw objects list (useful for visualization).

Target Encoding: Implementing Responsibility

The `encode_target` method is where YOLO's grid-based formulation becomes concrete. It creates a tensor [S, S, B, 5+C] initialized to zeros. For each object, we first determine which grid cell contains its center: $\text{grid_x} = \text{int}(x \times S)$, $\text{grid_y} = \text{int}(y \times S)$. This is the responsibility assignment - the cell containing the center is responsible for detecting the object.

We compute cell-relative coordinates: $x_{\text{cell}} = x \times S - \text{grid_x}$, $y_{\text{cell}} = y \times S - \text{grid_y}$. These are in [0, 1] range within the cell. For example, if an object's center is at $x=0.37$ and $S=7$,

`grid_x=2` (it's in the 3rd cell horizontally), and $x_cell = 0.37 \times 7 - 2 = 0.59$ (it's 59% across that cell).

The implementation simplifies responsibility by always assigning to box index 0. In the original YOLO, you'd compute IoU between the object's ground truth box and both predicted boxes, assigning responsibility to whichever has higher IoU. For synthetic data with well-separated objects, this simplification works fine.

The target tensor at `[grid_y, grid_x, box_idx]` is filled with: `x_cell`, `y_cell` (offsets within cell), `w`, `h` (absolute dimensions), `confidence=1.0` (object present), and one-hot class encoding. This structure mirrors the model's output, making the loss computation straightforward.

The Training Function: Bringing It All Together

The `train_yolo` function orchestrates the complete training process. It sets up the dataset, model, loss, and optimizer, then runs the training loop with detailed monitoring.

The dataset is instantiated with 1000 synthetic samples, providing sufficient diversity for learning the basic detection task. The DataLoader batches these with shuffling, ensuring varied training batches. The model is created with matching parameters (3 classes, 7×7 grid, 2 boxes per cell) and moved to GPU if available.

The loss function gets the same parameters as the model - they must match for the shapes to align. The Adam optimizer with learning rate $1e-3$ provides adaptive learning rates, typically converging faster than SGD for this task.

The training loop runs for the specified epochs, maintaining a history dictionary to track all loss components. For each batch, we move data to device, run the forward pass to get predictions, compute the multi-component loss, backpropagate, and update weights. The loss dictionary from YOLOLoss lets us track each component separately.

The epoch summary computes averages and prints all components. This detailed logging is educational - students can see how each loss term evolves. Typically, localization losses decrease steadily as the network learns where objects are, confidence losses balance as the network learns when to be confident, and classification loss drops as class discrimination improves.

The return values (model, history, dataset) allow for further analysis and visualization. The trained model can be evaluated, the history can be plotted, and the dataset is needed for visualization.

The Visualization Functions: Understanding Predictions

The `visualize_predictions` function demonstrates the trained model's performance. It sets the model to eval mode, processes several test samples, and visualizes both ground truth and predictions side by side.

For each sample, we get the image, target, and ground truth objects. The prediction is obtained by passing the image through the model (with a batch dimension added and removed). The `decode_predictions` function converts the raw YOLO output into human-readable detections.

The plotting draws ground truth boxes with dashed lines and prediction boxes with solid lines, both colored by class. Text annotations identify each box - "GT: class_name" for ground truth, "class_name confidence" for predictions. This visual comparison immediately shows where the model succeeds and struggles.

Decoding Predictions: From Grid to Boxes

The `decode_predictions` function converts YOLO's structured output [S, S, B, 5+C] into a list of detections. This is the inverse of encoding - we're going from grid-based representations back to bounding boxes in image coordinates.

For each grid cell (i, j) and each box (b), we extract the prediction components: cell-relative coordinates (x_{cell} , y_{cell}), dimensions (w, h), confidence, and class probabilities. The class with highest probability is selected. The score is $\text{confidence} \times \text{class_probability}$ - this combines "is there an object?" with "what class is it?", giving an overall detection confidence.

Detections below the threshold are filtered out. For remaining detections, we convert to absolute coordinates: $x = (j + x_{cell}) / S$, $y = (i + y_{cell}) / S$. This maps the cell index and offset back to normalized image coordinates. The width and height are already absolute (they were predicted as image-relative dimensions), so we just take their absolute value.

The result is a list of detections, each with x, y, w, h in normalized coordinates, plus confidence and class. This format is easy to visualize and evaluate.

Training Curves: Monitoring Convergence

The `plot_training_curves` function creates two panels showing loss evolution. The total loss panel shows overall convergence - a decreasing curve indicates successful learning. The components panel breaks this down, revealing which aspects of detection are improving.

Typically, you'd see different convergence patterns: localization losses decrease steadily and smoothly as the network learns spatial reasoning. Confidence losses often fluctuate more - balancing object and no-object confidence requires finding the right operating point. Classification loss usually decreases rapidly at first (learning to distinguish basic shape features) then plateaus (refining subtle differences).

These curves are diagnostic tools. If localization loss stays high, the network struggles with spatial reasoning. If confidence(noobj) dominates, the network might be predicting everything as background. If classification loss is stuck, the network can't distinguish classes. Understanding these patterns helps debug training issues.

The Historical Significance: YOLO's Impact

Understanding YOLO's implementation reveals why it was transformative. Before YOLO, detection was inherently multi-stage: generate proposals, extract features, classify. This pipeline was computationally expensive and indirect. YOLO proved you could reformulate detection as a single regression problem - directly predict all boxes and classes in one forward pass.

The grid-based formulation was the key insight. By dividing the image into cells and having each cell predict boxes, YOLO naturally handles multiple objects while maintaining spatial coherence. The responsibility concept (the cell containing an object's center is responsible for detecting it) provides a clear training signal without complex assignment algorithms.

YOLO's speed was revolutionary: 45 FPS on a GPU, enabling real-time detection for the first time. This wasn't just a theoretical achievement - it enabled new applications in robotics, autonomous vehicles, and video analysis where fast detection is critical. The trade-off was accuracy - YOLO v1 lagged behind Faster R-CNN, especially for small objects and precise localization.

But YOLO established a paradigm. The one-stage, direct prediction approach inspired SSD, which added multi-scale feature maps. RetinaNet proved one-stage could match two-stage accuracy with the right loss. Modern YOLO versions (v3, v4, v5, v8) incorporate ideas from across detection history while maintaining the core grid-based philosophy. Even transformer-based detectors like DETR can be seen as extensions of YOLO's direct prediction paradigm.

The Training Details: What Makes YOLO Work

The loss function's complexity is necessary because YOLO predicts everything simultaneously. The lambda weightings (5.0 for localization, 0.5 for no-object confidence) aren't arbitrary - they're carefully tuned to balance gradients across the different tasks. Without these, the overwhelming number of no-object cells would cause the network to predict low confidence everywhere, never learning to detect objects.

The square root for width/height is a subtle but important detail. It makes the loss more invariant to object size, preventing large objects from dominating the gradient. This design choice shows deep understanding of the regression problem - naive L2 loss on box dimensions would penalize errors on large boxes too heavily.

The cell-relative x, y encoding is crucial for learnability. If we tried to directly predict absolute image coordinates, different grid cells would need to learn vastly different mappings (top-left cells predict small values, bottom-right cells predict large values). With cell-relative coordinates, every cell learns the same task: given an object center in my cell, predict the offset within [0, 1].

The confidence definition (probability of object \times IoU with ground truth) elegantly combines objectness and localization quality in a single value. During inference, multiplying confidence by class probability gives a score that reflects "is there an object?", "is it localized well?", and "what class is it?" all at once.

The Pedagogical Value: Understanding Through Implementation

This YOLO implementation is pedagogically excellent because it demonstrates the complete pipeline with manageable complexity. The synthetic dataset removes real-world messiness, letting students focus on the algorithm. The detailed loss monitoring shows what each component does. The visualization makes predictions concrete.

Students learn several crucial concepts: grid-based detection (spatial structure in predictions), responsibility assignment (connecting objects to grid cells), multi-task learning (balancing different objectives), and direct prediction (bypassing proposals). These concepts extend beyond YOLO - they appear in semantic segmentation (grid-based pixel prediction), pose estimation (heatmap-based keypoint detection), and modern transformers (direct set prediction).

The contrast with region-based methods is also instructive. R-CNN and its descendants decompose detection into separate stages: propose, extract, classify. YOLO unifies these into a single network that predicts everything jointly. Neither approach is universally better - two-stage detectors typically achieve higher accuracy, one-stage detectors are faster. Understanding both paradigms gives students the full picture of detection strategies.

The implementation details matter too. The activation choices (sigmoid for offsets, softmax for classes), the loss weightings, the responsibility assignment - these aren't arbitrary. They're carefully designed solutions to specific problems that arise when trying to train a network to predict everything simultaneously. Studying this code teaches engineering judgment: how to design losses, balance objectives, and structure predictions for effective learning.

This is why we study YOLO deeply: it represents a fundamentally different approach to detection, one that proved you don't need explicit proposals or multi-stage processing. The grid-based, direct prediction paradigm it introduced influences modern detection to this day. Understanding this code means understanding how to rethink computer vision problems, moving from pipelines to end-to-end learning.

Intuitive Code Walkthrough for Lecture 6: Object Detection Methods

SSD (Single Shot MultiBox Detector) - Code Walkthrough

Let me walk you through how SSD actually works by looking at our implementation. Now, remember from our earlier discussion that by the time we get to SSD, we've moved away from the two-stage pipeline of R-CNN and Faster R-CNN. We're now in the realm of single-shot detectors, where we make all our predictions in one forward pass. This is the key philosophical shift that makes SSD fast enough for real-time applications.

The journey starts with our backbone network, which is really just a simple CNN that progressively downsamples the input image. What's interesting here is that unlike YOLO, which only uses the final feature map, SSD extracts features at multiple scales. In our implementation, we have what we call P3 and P4 levels, operating at strides of 16 and 32 respectively. Think about what this means: at P3, each spatial location in our feature map corresponds to a 16×16 pixel region in the original image, while at P4, each location represents a 32×32 region. This multi-scale approach is crucial because objects in natural images come in vastly different sizes. Small objects need the finer-grained P3 features where we still have good spatial resolution, while large objects can be effectively detected at P4 where we have more semantic, high-level features.

Now, here's where SSD introduces something really elegant: the anchor boxes, which we also call default boxes. At each location in each feature map, we don't just predict one bounding box. Instead, we predict multiple boxes with different aspect ratios and scales. Our implementation uses three aspect ratios (0.5, 1.0, and 2.0) and two scales (1.0 and 1.6), giving us six anchor boxes at each spatial location. Why do this? Well, think about the diversity of objects in the real world. Some are tall and thin like a person standing up, others are wide and short like a car, and some are roughly square. By having anchors with different shapes pre-defined, we're essentially giving the network templates that it can adjust, rather than trying to predict box dimensions from scratch. This makes the learning problem much easier.

The matching process is where things get really interesting from a training perspective. During training, we need to decide which anchor boxes are responsible for detecting which ground truth objects. We use Intersection over Union, or IoU, as our metric. If an anchor has IoU greater than or equal to 0.5 with a ground truth box, we call it a positive match and train it to refine its prediction toward that ground truth. If the IoU is less than 0.4, it's a negative, meaning there's no object there, and we train it to predict background. The anchors that fall between 0.4 and 0.5? We simply ignore them during training. This ignore zone is actually quite clever because it prevents the network from getting confused by anchors that are sort of close to an object but not really responsible for it.

But here's a problem: in a typical image, we might have thousands of anchors but only a handful of actual objects. This creates a massive class imbalance where negatives vastly outnumber positives. SSD's solution is something called hard negative mining. Instead of using all the negative anchors, we only keep the ones the network finds most confusing, the ones with the highest classification loss. Specifically, we maintain a 3:1 ratio of negatives to positives. So if we have 10 positive anchors, we only use the 30 hardest negative examples. This forces the network to focus on the difficult cases rather than getting swamped by easy background patches.

The loss function in SSD is a multi-task loss that jointly optimizes classification and localization. For classification, we use standard cross-entropy loss, but only on the selected positives and hard negatives. For localization, we use smooth L1 loss, which you can think of as a robust version of mean squared error that doesn't penalize outliers as harshly. Critically, we only compute localization loss on positive anchors because it doesn't make sense to regress a bounding box when there's no object there. The predictions themselves are treated as residuals or offsets from the anchor positions, so the network is learning to adjust the anchors rather than predict absolute coordinates.

When we decode predictions at inference time, we first apply sigmoid to get class probabilities, then we threshold them to keep only confident predictions. But we're not done yet because we'll have many overlapping detections for the same object, especially from different pyramid levels or neighboring anchors. This is where Non-Maximum Suppression comes in. We do NMS independently for each class: we sort all detections by confidence, keep the highest one, and then remove any other detections that have high IoU overlap with it. We repeat this process until we've either kept or suppressed every detection. What emerges from NMS is a clean set of final predictions, one box per object.

RetinaNet - Code Walkthrough

Now let's talk about RetinaNet, which represents the next evolution in single-shot detection. RetinaNet was born from a simple question: can we make one-stage detectors as accurate as two-stage detectors? The answer turned out to be yes, but it required solving a fundamental problem that had been holding back single-shot methods.

The architecture starts with something called a Feature Pyramid Network, or FPN, and this is more sophisticated than SSD's simple multi-scale approach. In SSD, we just take features at different depths of the network. In RetinaNet's FPN, we do something smarter. We have a bottom-up pathway, which is just our regular convolutional backbone creating features at progressively lower resolutions. But then we add a top-down pathway that takes these high-level semantic features and enriches the lower-level features with them. Concretely, we take P4 features, which are semantically rich but spatially coarse, upsample them, and add them to P3 features through lateral connections. This gives us the best of both worlds: P3 now has both fine spatial detail and high-level semantic information. Think of it as the network asking, "I see edges and textures at P3, but what are they part of?" and the top-down pathway answering, "based on P4, they're part of a person."

The prediction heads in RetinaNet are also more sophisticated than SSD's. Instead of a single convolutional layer, we have dedicated subnets with four convolutional layers each, one for classification and one for box regression. These are shared across all pyramid levels, which means the same weights process P3 and P4. This sharing has a regularization effect and reduces the number of parameters, but it also means the network learns scale-invariant feature representations. The subnet architecture is deeper because the features coming from FPN are richer and can support more complex prediction functions.

Now we come to the key innovation: focal loss. Remember how SSD dealt with class imbalance through hard negative mining, manually selecting which negatives to use? Focal loss handles this automatically through a clever reformulation of the cross-entropy loss. The

standard cross-entropy loss treats all examples equally, so even easy negatives where the network is already very confident contribute to the gradient. Focal loss adds a modulating factor: $(1 - p_t)^\gamma$. When the network is very confident and correct (p_t is close to 1), this term becomes nearly zero, effectively removing easy examples from the loss. When the network is uncertain or wrong (p_t is small), the term stays close to 1, so these hard examples still contribute fully. The gamma parameter, which we set to 2, controls how aggressively we down-weight easy examples.

There's also an alpha term in focal loss that provides class balancing, similar to weighted cross-entropy. We set alpha to 0.25 for positive classes, which means negatives get a weight of 0.75. But the real magic is in that focusing parameter. With focal loss, we don't need hard negative mining at all. We can use every single anchor during training, and the loss function automatically focuses the learning on the examples that matter. This is not just elegant from an algorithmic perspective; it actually works better because we're not throwing away information by discarding negatives.

Another subtle but important difference is that RetinaNet doesn't have an explicit background class. Instead, it uses sigmoid activation and treats each class independently. During training, if an anchor is negative, all its class targets are zero. If it's positive, only the target class has a 1, and others remain 0. This is different from SSD's softmax approach where we have a dedicated background class. The sigmoid approach is more flexible because it naturally handles multi-label scenarios, though in our simple shapes dataset, each object has only one class anyway.

The initialization of the final classification layer is also worth noting. We set the bias to a special value that makes the network initially predict very low probabilities for all classes. Why? Because at the start of training, almost all anchors are negatives, so if we initialized randomly, we'd get many false positives with high confidence. The focal loss would down-weight them as easy examples, and the network might never learn properly. By starting with low confidences, we ensure the network gradually learns to make positive predictions only when it's seen enough evidence.

During inference, the decoding process is similar to SSD in spirit but differs in the details. We apply sigmoid to get independent class scores, threshold them, and apply NMS per class. But because we're using sigmoid instead of softmax, we could theoretically detect multiple classes at the same location, though this rarely happens in practice. The NMS procedure is the same: sort by confidence, keep the best, suppress overlapping detections, repeat.

What's beautiful about comparing SSD and RetinaNet is that they solve the same fundamental problem—class imbalance in single-shot detection—but in completely different ways. SSD uses hard negative mining, a discrete selection mechanism that requires careful tuning of the negative-to-positive ratio. RetinaNet uses focal loss, a continuous, differentiable solution that automatically adjusts the contribution of each example based on how hard it is. Both work, but focal loss has proven to be more robust and has become the standard approach in modern object detection. The FPN architecture has also become ubiquitous because that combination of spatial detail and semantic richness is exactly what we need for detecting objects at multiple scales.

As you can see, the evolution from YOLO to SSD to RetinaNet represents progressively more sophisticated solutions to the core challenges of single-shot detection: handling multiple scales, dealing with class imbalance, and making accurate predictions in one forward pass. Each method builds on the insights of its predecessors while introducing new ideas that push the field forward.

A Detailed Class-by-Class Walkthrough of FCOS

Let me guide you through FCOS (Fully Convolutional One-Stage Object Detection), which represents a fundamental shift in how we think about object detection. While previous methods - from R-CNN to RetinaNet to YOLO - all rely on predefined anchor boxes or grid-based proposals, FCOS asked a radical question: what if we could detect objects without any anchors at all? What if every pixel in the feature map could directly predict whether it's inside an object and how far it is from that object's boundaries?

The Anchor-Free Revolution: Why FCOS Matters

Before diving into the code, understand the philosophical shift FCOS represents. Anchor-based methods (Faster R-CNN, RetinaNet) require carefully designed anchor boxes at multiple scales and aspect ratios - you need to predict which anchor is closest to your object, then refine it. YOLO uses a grid-based approach where cells predict boxes, which is conceptually similar to anchors. All these methods share a common assumption: we need predefined templates that we classify and adjust.

FCOS challenged this assumption. The insight was elegantly simple: in semantic segmentation, every pixel predicts a class without any templates. Why can't we do the same for detection? FCOS treats detection as per-pixel classification (is this pixel inside an object?) plus per-pixel regression (how far is this pixel from the object's boundaries?). This formulation is anchor-free, simpler to implement, and eliminates all anchor-related hyperparameters (scales, aspect ratios, IoU thresholds for anchor assignment).

The key innovation enabling this is the **Itrb formulation**: for each pixel at location (x, y) , instead of predicting a box directly, we predict four distances - left, top, right, bottom - to the object's boundaries. If a pixel is inside an object, these four distances uniquely define the bounding box. This representation is natural (distances from a point to boundaries), scale-equivariant (nearby pixels predict similar distances), and handles arbitrary shapes.

But there's a problem: pixels far from the object center give poor bounding boxes (if you're at the very edge of an object, the "left" distance is nearly zero while the "right" distance spans the entire object - this asymmetry produces low-quality boxes). FCOS solves this with the **centerness branch**: a parallel prediction that estimates how close a pixel is to the object

center. During inference, we multiply the classification score by centerness, effectively suppressing low-quality predictions from pixels near object boundaries.

The SimpleBackbone Class: Foundation for Feature Extraction

The `SimpleBackbone` class implements a ResNet-like architecture that extracts multi-scale features for FCOS. While FCOS can use any backbone, the key requirement is producing features at multiple resolutions that will feed into the FPN.

The initialization builds a hierarchical network with clear downsampling stages. The first convolution (`conv1`) uses a large 7×7 kernel with stride 2, immediately reducing spatial dimensions to half while increasing receptive field. This is followed by batch normalization, ReLU activation, and max pooling with stride 2, giving us an initial downsampling factor of 4. This aggressive early downsampling is standard in modern architectures - it quickly reduces computation while building up receptive field.

The subsequent layers (`layer1` through `layer4`) progressively downsample and increase channels. The `_make_layer` helper constructs blocks of convolutions - the first conv in each layer handles downsampling (via `stride=2` where needed) and channel expansion, while subsequent convs maintain dimensions and provide additional capacity. The channel progression ($64 \rightarrow 128 \rightarrow 256 \rightarrow 512$) and stride progression ($4 \rightarrow 8 \rightarrow 16 \rightarrow 32$) create a feature hierarchy suitable for multi-scale detection.

The `forward` method extracts features at multiple stages and returns them as a dictionary. Critically, it returns `C3`, `C4`, and `C5` - features at strides 8, 16, and 32 respectively. These correspond to different receptive field sizes and semantic levels: `C3` captures fine details with smaller receptive fields, `C5` captures high-level semantics with large receptive fields. The FPN will fuse these to create features that have both spatial precision and semantic understanding at each scale.

The FPN Class: Multi-Scale Feature Fusion

The `FPN` class implements the Feature Pyramid Network we've seen before, but its role in FCOS is worth emphasizing. FCOS is fundamentally a per-pixel prediction method, so having high-quality features at each scale is crucial - poor features lead to poor predictions at every pixel.

The initialization sets up lateral connections and output convolutions just as before: 1×1 convs to align channels, 3×3 convs to smooth the fused features. The key architectural choice is the unified output channels (256) across all pyramid levels - this allows FCOS to use a single shared detection head for all scales.

The `forward` method builds the pyramid through the top-down pathway: start with `P5` (deepest features), upsample and add to `C4` to create `P4`, upsample `P4` and add to `C3` to create `P3`. Each addition fuses information: the lateral connection preserves spatial precision

from that level, the upsampled features bring semantic understanding from deeper levels. The output convolutions smooth the result.

For FCOS, this multi-scale architecture is essential because objects of different sizes are detected at different pyramid levels. Small objects (8-64 pixels) are detected on P3 with its high resolution (28×28 for 224×224 input). Medium objects (64-128 pixels) use P4's 14×14 resolution. Large objects (128+ pixels) use P5's 7×7 resolution. This scale assignment is automatic - it depends on which pyramid level a pixel belongs to and what object it's inside.

The Scale Class: Learnable Scaling for Regression

The `Scale` class is a beautifully simple component that addresses a practical problem in multi-scale detection. It's a single learnable scalar parameter that multiplies the regression outputs.

The initialization creates a scalar parameter initialized to 1.0. This parameter is learned during training via backpropagation. The `forward` method simply multiplies the input by this scalar.

Why is this needed? Different pyramid levels detect objects of vastly different sizes. P3 might predict distances of 10-30 pixels (small objects), while P5 predicts distances of 100-300 pixels (large objects). If we use the same regression head for all levels (which FCOS does for efficiency), the network must learn to produce appropriate magnitudes at each level. The learnable scale provides a per-level correction factor: the network can learn to produce normalized outputs, and the scale adjusts them to the appropriate magnitude for each pyramid level.

In practice, the scales typically evolve to be larger for deeper levels (P5 gets a larger scale than P3), reflecting the larger objects those levels detect. This is a clever way to handle multi-scale regression with a shared head - the head produces normalized predictions, and the scales provide level-specific calibration.

The FCOSHead Class: Per-Pixel Predictions

The `FCOSHead` class is where FCOS's anchor-free magic happens. This head makes predictions at every spatial location: classification (what class?), regression (where are the boundaries?), and centerness (how good is this prediction?).

The initialization constructs three parallel towers, each with multiple convolutional layers. The classification tower processes features through four 3×3 convs, each followed by GroupNorm and ReLU. GroupNorm is chosen over BatchNorm because it's more stable for small batch sizes common in detection. The tower uses the same channel dimension throughout (256 by default), providing substantial representational capacity. The final layer is a 3×3 conv that outputs `num_classes` channels - one logit per class per spatial location.

The regression tower has an identical structure but outputs 4 channels representing the `ltrb` distances. This parallel design - separate towers for classification and regression - follows

RetinaNet's architecture and reflects the insight that these tasks benefit from specialized feature processing.

The centerness branch is simpler: it shares the regression tower's features (they're closely related - both concern box geometry) and adds a single 3×3 conv outputting 1 channel. This binary score estimates how close a location is to the object center.

The scales are stored as a `ModuleList` - one `Scale` module per pyramid level. When processing features from P3, we use `scales[0]`; for P4, `scales[1]`; and so on. This per-level scaling handles the multi-scale regression issue discussed earlier.

Weight Initialization: Setting the Right Starting Point

The `_init_weights` method includes crucial initialization choices. The tower convolutions use small random weights ($\text{std}=0.01$) and zero biases - standard initialization for detection networks. But the classification output gets special treatment: its bias is initialized to -2.19, which is $-\log((1-\pi)/\pi)$ for $\pi=0.01$.

This initialization is critical for training with focal loss. At the start of training with random weights, the network outputs roughly zero logits. Applying sigmoid gives 0.5 probability for all classes at all locations. But with massive class imbalance (most pixels aren't inside objects), this is wildly wrong and produces huge loss. By biasing the output to give initial probabilities around 0.01 (matching the actual object frequency), we start training from a reasonable point.

The regression and centerness heads get standard zero-bias initialization. For regression, the ReLU activation ensures positive outputs (distances must be positive). For centerness, we'll apply sigmoid during forward pass to get a probability-like score.

The Forward Pass: Making Predictions

The `forward` method processes features from a single FPN level and produces three outputs: classification logits, regression predictions, and centerness scores. It receives features $[B, 256, H, W]$ and a level index (0 for P3, 1 for P4, 2 for P5).

The classification branch passes features through the `cls_tower`, producing intermediate features $[B, 256, H, W]$, then through `cls_logits` to get $[B, \text{num_classes}, H, W]$. Each spatial location has `num_classes` logits.

The regression branch similarly processes through `reg_tower`, then through `reg_pred` to get $[B, 4, H, W]$ raw predictions. These are scaled by the level-specific `Scale` module - this is where per-level calibration happens. Finally, ReLU ensures all distances are positive (you can't have negative distance to a boundary). The output is $[B, 4, H, W]$ representing `ltrb` distances at each location.

The centerness branch shares the regression features (they're related to box geometry) and outputs $[B, 1, H, W]$. No activation is applied here - sigmoid will be used during loss computation and inference.

The returned tuple (cls_logits, reg_pred, centerness) provides complete predictions for this FPN level. The same head processes all levels, but the level_idx determines which scale is applied. This shared-head design is efficient and provides useful inductive bias - the same visual patterns indicate objects at all scales.

The FCOSDetector Class: Complete Architecture

The `FCOSDetector` class assembles the complete FCOS pipeline: backbone, FPN, and detection head. Understanding this class reveals how FCOS achieves anchor-free detection at multiple scales.

The initialization stores fundamental parameters: number of classes and the strides of the FPN levels [8, 16, 32]. These strides are critical because they define the relationship between feature map locations and image coordinates. A location (i, j) on P3 (stride 8) corresponds to pixel ($i \times 8 + 4, j \times 8 + 4$) in the image.

The backbone and FPN are instantiated with appropriate channels. The backbone outputs C3, C4, C5 with 128, 256, 512 channels respectively. The FPN unifies these to 256 channels, producing P3, P4, P5. The FCOS head is created with 256 input channels and `num_classes` outputs.

The crucial design choice is using a single shared head for all FPN levels. This is efficient (fewer parameters) and provides inductive bias (objects look similar at all scales). The level-specific scales provide the necessary per-level adaptation.

The Forward Pass: Multi-Scale Predictions

The `forward` method orchestrates the complete detection pipeline. Images [B, 3, H, W] flow through the backbone, producing the multi-scale feature dictionary. The FPN fuses these into the pyramid [P3, P4, P5], each with shape [B, 256, H_level, W_level].

For each FPN level, we apply the shared head with the appropriate level index. This produces `cls_logits` [B, num_classes, H, W], `reg_pred` [B, 4, H, W], and `centerness` [B, 1, H, W] for that level. These are collected into lists, one entry per FPN level.

The return values are three lists: `cls_logits_list`, `reg_preds_list`, and `centerness_list`. Each list has three elements (one per FPN level), and each element is a tensor with predictions for all spatial locations at that level. This structure preserves the multi-scale organization - we know which predictions came from which scale, which matters for both training (assigning ground truth to the right scale) and inference (interpreting box sizes).

The FocalLoss Class: Handling Class Imbalance

The `FocalLoss` class implements focal loss, which we've seen in RetinaNet, but its role in FCOS is worth emphasizing. FCOS makes predictions at every pixel, leading to extreme class imbalance - the vast majority of pixels aren't inside any object.

The implementation is the same as before: compute probability p_t for the correct class, apply modulating factor $(1-p_t)^\gamma$ to down-weight easy examples, and use alpha weighting for class balance. For FCOS, this is even more critical than RetinaNet because the prediction density is higher - every pixel predicts, not just anchor centers.

The forward method computes focal loss element-wise and sums. The sum (not mean) is used because we'll normalize by the number of positive samples in the main loss function. This normalization is important - if we average over all pixels, the loss would be dominated by the overwhelming number of negatives even with focal loss's down-weighting.

The IoULoss Class: Better Localization

The `IoULoss` class implements IoU (Intersection over Union) loss for bounding box regression, which is superior to L1 or L2 loss for boxes. The insight is that we should directly optimize the metric we care about (IoU) rather than a proxy (coordinate errors).

The forward method receives predicted and target l_{trb} distances [$N, 4$] where N is the number of positive pixels. For each pixel, we have predicted (l, t, r, b) and target (l^*, t^*, r^*, b^*) distances. The implementation computes IoU between the boxes defined by these distances.

The box reconstruction is straightforward: the box width is $l + r$ (left distance plus right distance), and height is $t + b$. The area is width \times height. For intersection, we take the minimum of corresponding distances: the intersection width is $\min(l, l^*) + \min(r, r^*)$, because the intersection can't extend further than the smaller box in any direction. The intersection area is `intersection_width` \times `intersection_height`.

Union area is sum of areas minus intersection (to avoid double-counting). IoU is intersection/union. The loss is $1 - \text{IoU}$, so perfect alignment gives loss 0, and no overlap gives loss 1.

Why is this better than L1 or L2 loss on coordinates? Consider two predictions: one has all distances off by 5 pixels on a 50×50 box (10% error), another has all distances off by 5 pixels on a 500×500 box (1% error). L1 loss treats these equally (same absolute error), but IoU loss correctly assigns higher loss to the first case (worse IoU). IoU loss is scale-invariant, matching how we evaluate detection quality.

The FCOSLoss Class: Multi-Component Training Objective

The `FCOSLoss` class combines focal loss, IoU loss, and centerness loss into the complete FCOS training objective. Understanding this class reveals how FCOS trains its per-pixel predictions effectively.

The initialization stores the number of classes and strides, and instantiates the focal loss and IoU loss modules. The strides aren't used directly in loss computation but are included for consistency with the architecture.

Computing Centerness Targets

The `compute_centerness_targets` method computes the ground truth centerness values from regression targets. The centerness definition is clever: for a pixel with distances (l, t, r, b) to an object's boundaries, centerness is $\sqrt{[(\min(l,r) / \max(l,r)) \times (\min(t,b) / \max(t,b))]}.$

This formula captures "how centered" a pixel is within the object. If a pixel is perfectly centered, $l \approx r$ and $t \approx b$, giving centerness ≈ 1 . If a pixel is near the left edge, $l \ll r$, giving $\min(l,r)/\max(l,r) \approx 0$, so centerness ≈ 0 . The square root provides a slight smoothing that works well empirically.

The implementation computes left-right and top-bottom min/max ratios separately, multiplies them, and takes the square root. This gives targets in $[0, 1]$ where 1 indicates perfect centering. During training, the network learns to predict high centerness for well-centered pixels and low centerness for edge pixels.

The Forward Pass: Multi-Level Loss Computation

The `forward` method computes the complete loss across all FPN levels. It receives lists of predictions (one per level) and a list of targets (one per level). For each level, predictions have shape $[B, C/4/1, H, W]$ which we reshape to $[B \times H \times W, C/4/1]$ to process all locations together.

The labels tensor indicates which pixels are positive (inside an object with class ≥ 0) or negative (background, class = -1). The positive mask selects pixels that should contribute to regression and centerness losses. The number of positives is computed with a minimum of 1 to avoid division by zero.

For classification loss, we convert labels to one-hot encoding, but set background pixels (negative mask) to all zeros. This makes them contribute to focal loss as negatives. The focal loss is normalized by num_pos - this is crucial because without normalization, levels with more positives would dominate the gradient.

For regression loss, we only consider positive pixels. Their predicted and target ltrb distances go into IoU loss, which computes $1 - \text{IoU}$ for each box and sums. Again, we normalize by num_pos. This normalization ensures that images with many objects don't dominate training.

For centerness loss, we first compute targets from the regression targets using our formula. Then we compute binary cross-entropy between predicted and target centerness, only for positive pixels. The BCE with logits applies sigmoid internally, so our raw centerness predictions (before sigmoid) are compared to the $[0,1]$ targets. Normalization by num_pos ensures consistent loss magnitudes.

The final aggregation averages losses across FPN levels. This gives equal weight to all scales, preventing any single scale from dominating. The total loss is simply the sum of the three components - no weighting is needed because they're already well-balanced in magnitude through the normalization.

The SyntheticObjectDataset Class: Adapted for FCOS

The `SyntheticObjectDataset` class generates training data in FCOS's format. While similar to YOLO's dataset, the target encoding is fundamentally different because FCOS is anchor-free and uses ltrb distances.

The initialization is similar to before: we set up synthetic shapes with known classes and colors. The key difference is that FCOS targets are per-level (we create separate targets for P3, P4, P5) and use ltrb format instead of bounding boxes.

The `__getitem__` method creates an image with random shapes, just as before. Objects are stored with corner coordinates (x_1, y_1, x_2, y_2) rather than center and dimensions, because ltrb distances are computed from corners. The `encode_targets` method does the FCOS-specific encoding.

Target Encoding: Per-Pixel Assignments

The `encode_targets` method creates targets for each FPN level. For each stride (8, 16, 32), we create a feature map of size $H \times W = \text{image_size}/\text{stride}$. At each location (i, j) , we compute the corresponding image coordinates: $\text{loc}_x = (j + 0.5) \times \text{stride}$, $\text{loc}_y = (i + 0.5) \times \text{stride}$. The 0.5 offset places us at the center of the pixel's receptive field.

For each location, we check if it's inside any object by testing if loc_x is between x_1 and x_2 , and loc_y is between y_1 and y_2 . If so, we compute ltrb distances: $l = \text{loc}_x - x_1$ (distance to left edge), $t = \text{loc}_y - y_1$ (distance to top edge), $r = x_2 - \text{loc}_x$ (distance to right edge), $b = y_2 - \text{loc}_y$ (distance to bottom edge).

These four distances uniquely define the bounding box from this location's perspective. We assign the object's class to `labels[i, j]` and store the ltrb distances in `reg_targets[i, j]`. Locations not inside any object are left at -1 (background).

This per-pixel assignment is FCOS's key advantage over anchor-based methods. There's no IoU computation, no anchor matching, no complex assignment rules. If a pixel is inside an object, it's responsible for predicting that object's boundaries. Period. This simplicity eliminates numerous hyperparameters and makes the training signal clearer.

The multi-level targets naturally handle scale: small objects appear on P3 (high resolution), large objects appear on P5 (low resolution). This isn't explicitly coded - it emerges from the fact that at coarser resolutions (larger strides), only larger objects occupy multiple pixels.

The Training Function: Anchor-Free Learning

The `train_fcos` function orchestrates FCOS training with its anchor-free formulation. The setup is similar to other detectors: create dataset, model, loss, and optimizer. The key difference is in how targets are organized.

The collate function batches data, but FCOS targets require special handling. Each sample has targets_list with three elements (P3, P4, P5 targets). When batching, we need to transpose from [batch][level] to [level][batch] so we can process each level's predictions and targets together. This transposition happens before the forward pass.

The training loop processes batches, moves data to device, transposes targets, and does forward-backward passes. The loss monitoring tracks all three components: focal loss for classification, IoU loss for regression, and BCE for centerness. Watching these losses reveals FCOS's training dynamics: focal loss decreases as the network learns where objects are, IoU loss improves as localization becomes more accurate, and centerness loss drops as the network learns to score predictions by quality.

The epoch summary prints all components, providing diagnostic information. If focal loss is stuck, the network struggles with classification. If IoU loss stays high, localization is poor. If centerness is high, the network isn't learning to distinguish centered from edge predictions. These insights guide debugging.

The Decoding Function: From Itrb to Boxes

The `decode_predictions` function converts FCOS's per-pixel predictions into final detections. This is where we see how FCOS's formulation comes together.

For each FPN level, we extract predictions: classification logits become probabilities via sigmoid, regression predictions are the Itrb distances, and centerness becomes a probability via sigmoid. At each location, we get the maximum classification score and its class. We multiply by centerness - this is FCOS's key quality control mechanism, down-weighting predictions from poorly-centered pixels.

Scores below threshold are filtered. For remaining detections, we decode the box. The location (i, j) in the feature map corresponds to image location ((j+0.5)×stride, (i+0.5)×stride). Given Itrb distances from this location, the box corners are: $x_1 = \text{loc}_x - l$, $y_1 = \text{loc}_y - t$, $x_2 = \text{loc}_x + r$, $y_2 = \text{loc}_y + b$. This reconstructs the bounding box from the per-pixel prediction.

The resulting detections include box coordinates, confidence ($\text{class_score} \times \text{centerness}$), and class. The centerness multiplication is crucial - it implements FCOS's strategy of having all pixels inside an object make predictions, but trusting centered pixels more than edge pixels. This is more robust than only using the center pixel (what if it's occluded or has ambiguous features?) while avoiding low-quality predictions from edge pixels.

The Historical Significance: Anchor-Free Detection

Understanding FCOS's implementation reveals why anchor-free detection matters. Anchor-based methods require designing anchor boxes - choosing scales, aspect ratios, and matching thresholds. These are hyperparameters that need tuning for each dataset. FCOS eliminates all of this. Every pixel simply predicts distances to boundaries if it's inside an object. No templates, no matching, no anchor-related hyperparameters.

The Itrb formulation is elegant because it's natural (distances from a point to boundaries) and scale-equivariant (nearby pixels predict similar values). Combined with FPN's multi-scale features, this handles objects of all sizes without explicit scale assignment - small objects naturally appear on high-resolution features, large objects on low-resolution features.

The centerness branch was the key to making this work. Early anchor-free methods struggled with ambiguous predictions from object edges. Centerness provides a learned quality score that identifies which predictions to trust. It's a simple idea (one additional conv layer and BCE loss) but critically important for anchor-free detection.

FCOS's impact was immediate. It proved anchor-free detection could match anchor-based accuracy while being simpler. This inspired follow-up work: ATSS reconsidered what "anchor-free" means, CornerNet detected corners instead of centers, CenterNet detected centers directly. Even DETR's transformer-based detection can be seen as an extension of anchor-free ideas - predicting boxes directly without templates.

The implementation teaches important lessons beyond FCOS itself. It shows that architectural assumptions (anchors are necessary) can be challenged. It demonstrates how per-pixel predictions work (critical for dense prediction tasks like segmentation). It reveals the importance of quality estimation (centerness) when making predictions at every location. And it illustrates how simple ideas (predict distances from pixels to boundaries) can enable major architectural changes.

This is why we study FCOS carefully: it represents a paradigm shift from anchor-based to anchor-free detection. Understanding this code means understanding how to rethink fundamental assumptions in computer vision, how to design training objectives for dense predictions, and how to build systems that are both simpler and more effective. The anchor-free revolution that FCOS helped start continues to influence modern detection - from FCOS itself to transformers like DETR, the idea of direct prediction without predefined templates has become central to modern vision.

DETR: Detection Meets Transformers

1 The big shift

Before DETR (Carion et al., 2020), every detector — Faster R-CNN, RetinaNet, YOLO, even FCOS — relied on:

- **local anchors / grids**, and
- **post-processing (NMS)** to merge overlapping boxes.

DETR asked:

“Can we detect objects **like a translation problem** — directly mapping an image to a *set* of boxes, end-to-end?”

So DETR replaced anchors and NMS with a **Transformer encoder–decoder** that performs *global reasoning*.

2Core intuition

- **Backbone CNN** extracts a feature map (spatial tokens).
- Add **2-D positional encodings** so the transformer knows where each patch came from.
- Flatten to a sequence → feed into the **Transformer encoder**.
 - Encoder performs global self-attention: each region “looks” at all others.
- The **decoder** starts with **learnable object queries** (e.g., 100 query vectors).
 - Each query competes to represent one object in the image.
 - Through cross-attention with the encoded features, it learns *what* and *where* that object is.
- The output of each query is passed to:
 - a **class head** → what object type it found,
 - a **box head** → normalized box coordinates ($[x_c, y_c, w, h]$).

So you get a fixed-size *set* of predictions (e.g. 100 boxes), regardless of how many objects actually exist.

3How it trains

Instead of anchor matching, DETR uses **Hungarian matching** to find the best one-to-one correspondence between predicted boxes and ground-truth boxes.

Loss =

[

```
\text{Classification (CE)} + \text{L1 box regression} + \text{GIoU loss.}  
]
```

Unmatched queries are trained as “*no-object*.”

This “set-based loss” means:

- Each object is predicted once.
- **No NMS** is needed — duplicates never arise.

4Why it's revolutionary

Old paradigm	DETR paradigm
Thousands of anchors	Fixed set of learned queries
Local receptive fields	Global attention context
Heuristic NMS	Optimal matching via Hungarian algorithm
Complex pipeline	End-to-end differentiable training

5Quick mental picture

The CNN extracts *what each patch sees* →
the transformer decides *which patches belong together* →
each learned query says “I claim this object” →
the decoder outputs the final box and label.

♦ 6In your demo code

- `SimpleBackbone` → produces feature map.
- `PositionalEncoding2D` → injects spatial coordinates.
- `TransformerEncoder/Decoder` → global reasoning + query decoding.
- `query_embed` → 100 learnable object tokens.
- `class_embed, bbox_embed` → prediction heads.
- `HungarianMatcher + DETRLoss` → optimal assignment & set loss.

Training = **image** → **transformer** → **100 boxes + labels** → **match** → **loss** → **backprop**.

7Takeaway line

- ◆ DETR is the first truly end-to-end object detector — no anchors, no proposals, no NMS — just pure attention and set prediction.